




CSCI 5902
Adv. Cloud Architecting
Fall, 2023
Instructor: Dr. Lu Yang

Week 5 – Lec 1
Kubernetes Foundations
Oct 6, 2023

Housekeeping and Feedback

- Start recording
- Start learning GCP CI/CD tools, cloud shell, and Terraform
- PIER tour spreadsheet available Monday afternoon, Oct 9

Oct 10




Part 1

Introduction to Kubernetes

Part 2

Kubernetes architecture



Part 1 Introduction to Kubernetes

- Containers on GCP
- Introduction to Kubernetes
- Introduction to Google Kubernetes Engine
- Google compute services
- Summary

Part 1 Introduction to Kubernetes

Introduction to Containers

How can you get or create containers?



Download containerized software from container registry or artifact registry such as gcr.io.

docker

Build your own container using the open-source docker command.



Build your own container using Cloud Build.

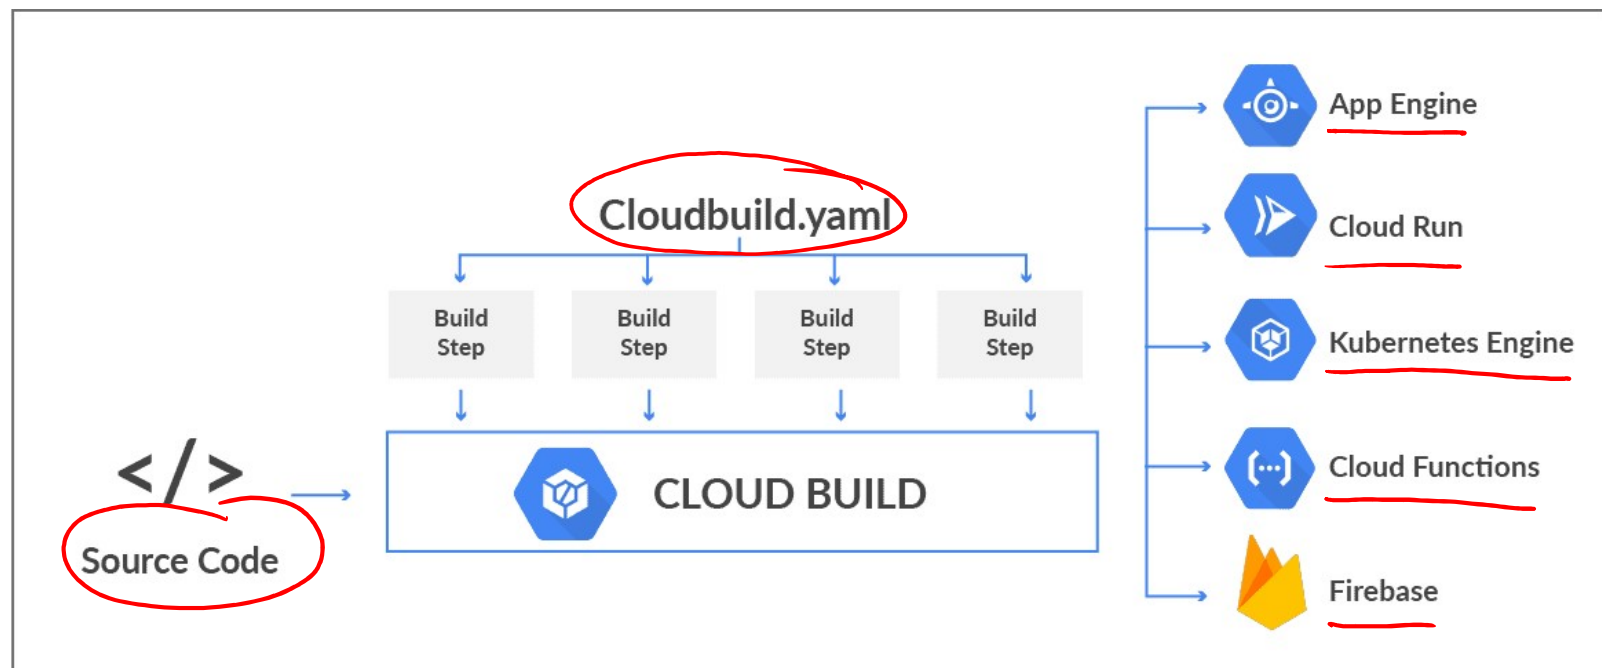
Artifact Registry is the recommended service for container image storage and management on Google Cloud. Artifact Registry provides the same container management features as Container Registry and includes additional features and benefits. As a fully-managed service with support for both container images and non-container artifacts, Artifact Registry extends the capabilities of Container Registry.


<https://cloud.google.com/artifact-registry/docs/transition/transition-from-gcr>

Part 1 Introduction to Kubernetes

Introduction to Containers

Build and deploy containers on GCP



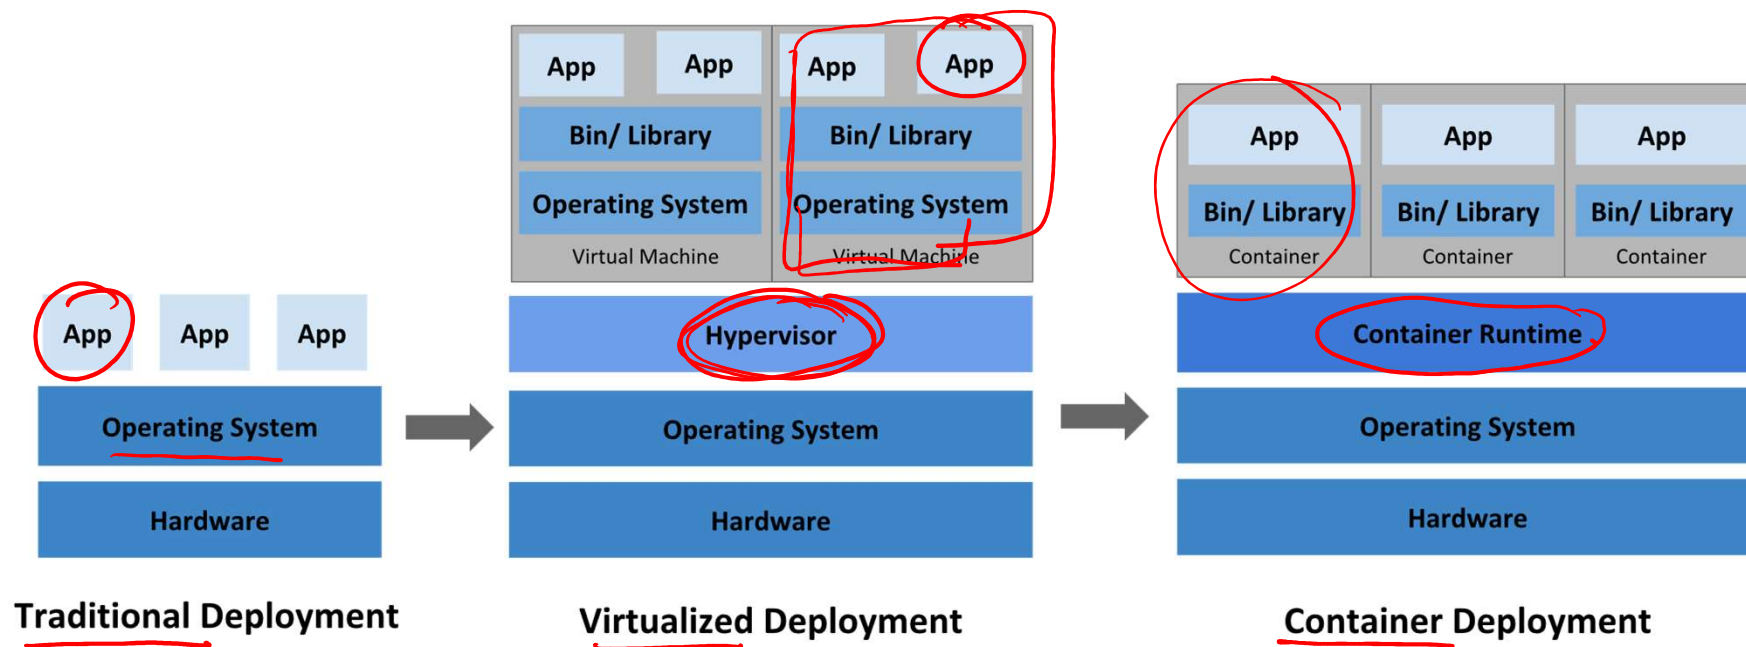


Part 1 Introduction to Kubernetes

- Introduction to Containers
- **Introduction to Kubernetes**
- Introduction to Google Kubernetes Engine
- Google compute services
- Summary

Part 1 Introduction to Kubernetes

Introduction to Kubernetes



How do you better manager your container infrastructure?

Kubernetes! (<https://kubernetes.io/docs/concepts/overview/>)

Part 1 Introduction to Kubernetes

Introduction to Kubernetes


What does Kubernetes do?


- **Service discovery and load balancing** Kubernetes can expose a container using the DNS name or using their own IP address. If traffic to a container is high, Kubernetes is able to load balance and distribute the network traffic so that the deployment is stable.
- **Storage orchestration** Kubernetes allows you to automatically mount a storage system of your choice, such as local storages, public cloud providers, and more.
- **Automated rollouts and rollbacks** You can describe the desired state for your deployed containers using Kubernetes, and it can change the actual state to the desired state at a controlled rate. For example, you can automate Kubernetes to create new containers for your deployment, remove existing containers and adopt all their resources to the new container.
- **Automatic bin packing** You provide Kubernetes with a cluster of nodes that it can use to run containerized tasks. You tell Kubernetes how much CPU and memory (RAM) each container needs. Kubernetes can fit containers onto your nodes to make the best use of your resources.
- **Self-healing** Kubernetes restarts containers that fail, replaces containers, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.
- **Secret and configuration management** Kubernetes lets you store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys. You can deploy and update secrets and application configuration without rebuilding your container images, and without exposing secrets in your stack configuration.

Part 1 Introduction to Kubernetes

Introduction to Kubernetes

What is Kubernetes?






Part 1 Intro to Kubernetes

- Introduction to Containers
- Introduction to Kubernetes
- Introduction to Google Kubernetes Engine
- Google compute services
- Summary

Part 1 Introduction to Kubernetes

Introduction to Google Kubernetes Engine



Google Kubernetes Engine

Part 1 Introduction to Kubernetes

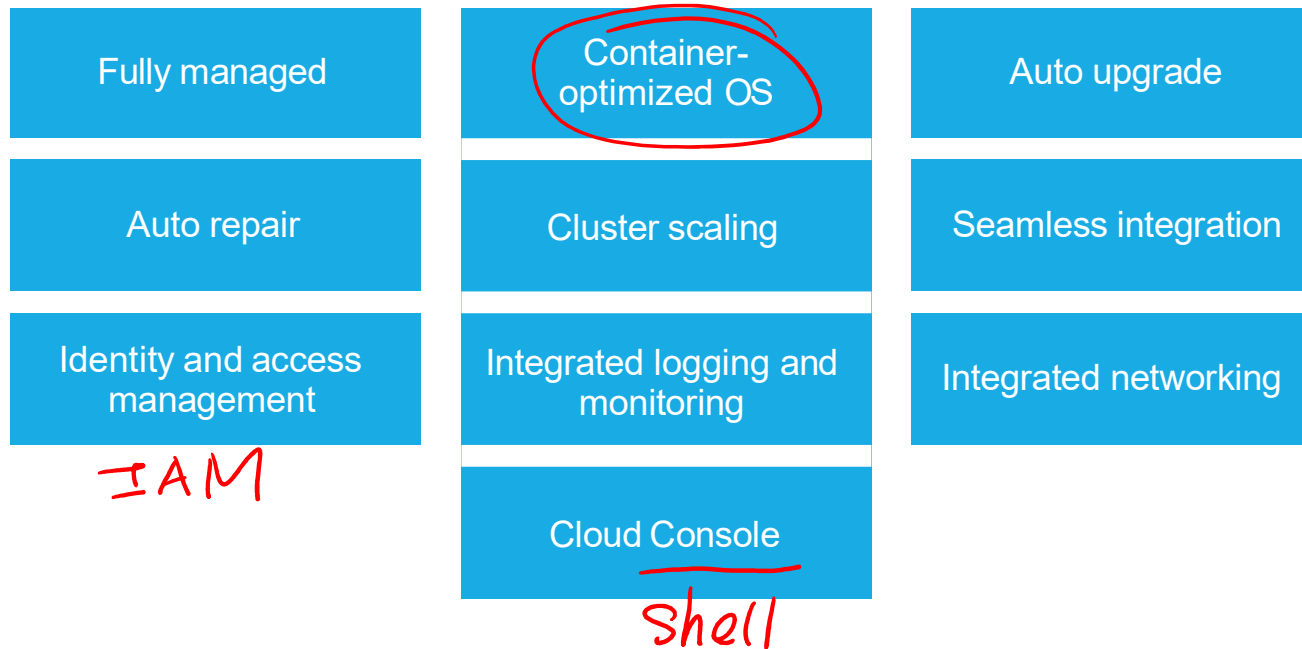
Introduction to Google Kubernetes Engine


- GKE is a managed Kubernetes service on Google infrastructure. GKE helps you to **deploy**, **manage**, and **scale** Kubernetes environments for your containerized applications on Google Cloud.
- More specifically, GKE is a component of the Google Cloud compute offerings. It makes it easy to bring your Kubernetes workloads into the cloud.

Part 1 Introduction to Kubernetes

Introduction to Google Kubernetes Engine

Features of GKE





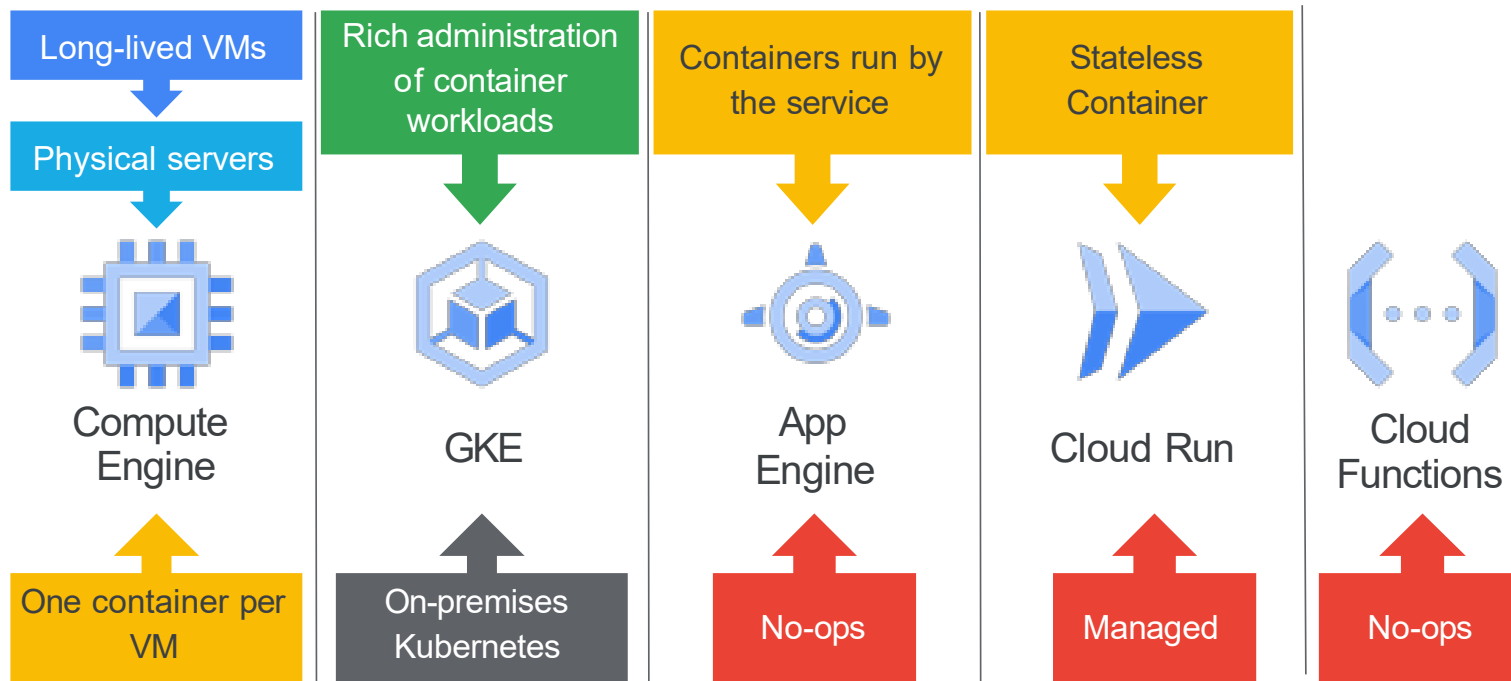
Part 1 Intro to Kubernetes

- Introduction to Containers
- Introduction to Kubernetes
- Introduction to Google Kubernetes Engine
- Google compute services
- Summary

Part 1 Introduction to Kubernetes

Google Compute Services

Which compute service should you adopt?

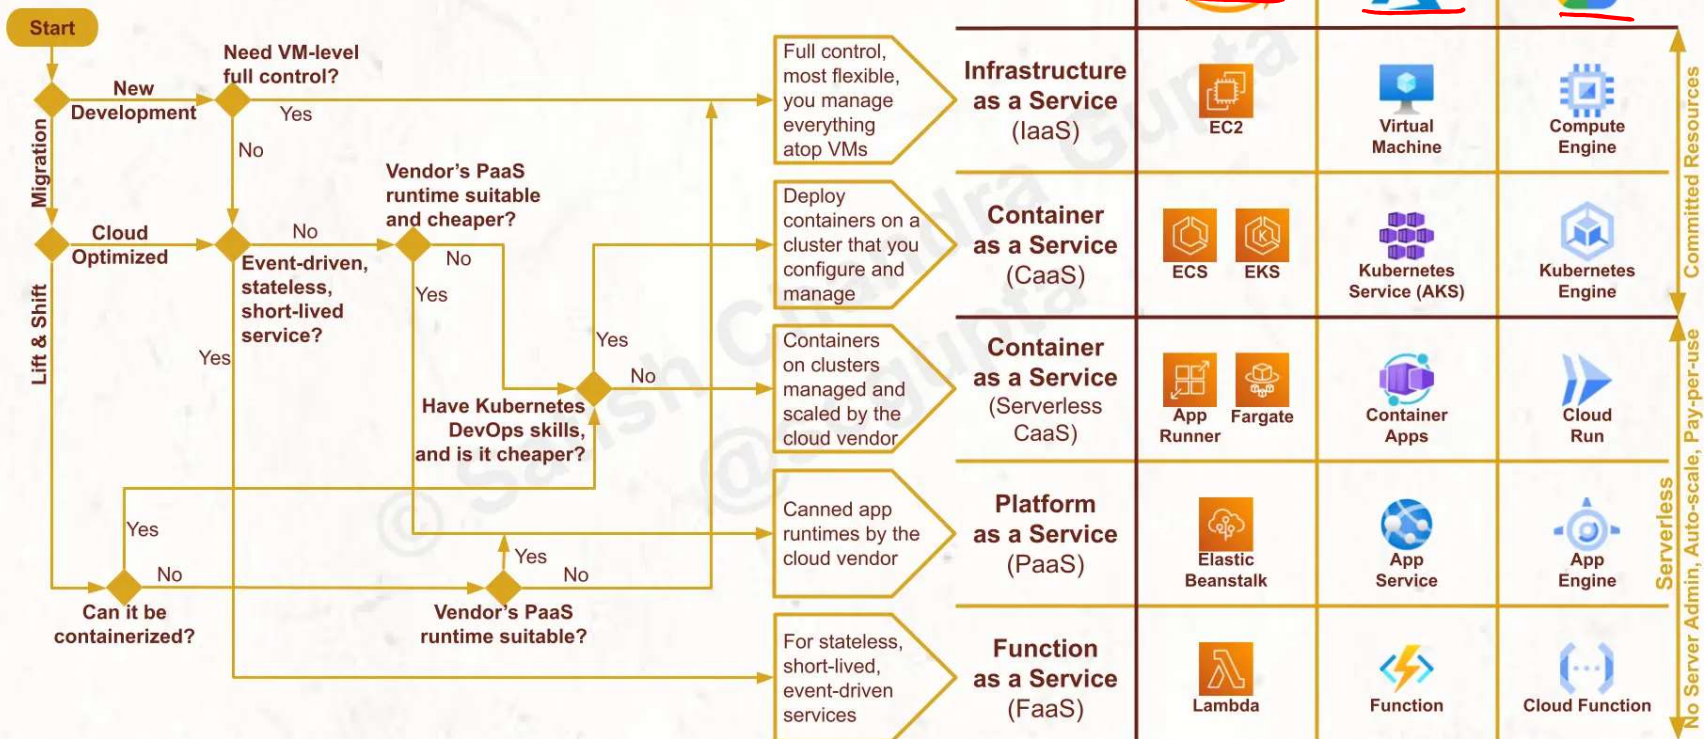


Part 1 Introduction to Kubernetes

Compute Services on AWS, Azure, and GCP

Cloud Deployment Cheatsheet for AWS, Azure, and Google Cloud

ml4devs.com/serverless



© Satish Chandra Gupta, Creative Commons BY-NC-ND 4.0 International License.

scgupta

linkedin.com/in/scgupta

<https://www.ml4devs.com/articles/serverless-architecture-for-microservices-on-aws-vs-google-cloud-vs-azure-as-iaas-caas-paas-faas/>

Part 1 Introduction to Kubernetes


Google Kubernetes Engine

Google Kubernetes Engine (GKE)

- Fully managed Kubernetes platform.
 - Supports cluster scaling, persistent disks, automated upgrades, and auto node repairs. *hard drive/EBS*
 - Built-in integration with Google Cloud services.
 - Hybrid computing
 - Multi-cloud computing
-

GKE use cases

- Containerized applications.
- Cloud-native distributed systems.
- Hybrid applications.



Part 1 Intro to Kubernetes

- Introduction to Containers
- Lab: Working with Cloud Build Introduction to Kubernetes
- Introduction to Kubernetes
- Introduction to Google Kubernetes Engine
- Computing Options
- Summary

Part 1 Introduction to Kubernetes

Summary


Summary

Create a container using Cloud Build. Google Cloud compute solutions:

- Compute Engine
- App Engine
- Google Kubernetes Engine
- Cloud Run
- Cloud Functions

Store a container in Artifact Registry.

Compare and contrast Kubernetes and Google Kubernetes Engine (GKE) features.



Part 2 Kubernetes Architecture

- **Kubernetes Concepts**
- GKE Architecture
- Object Management
- Lab: Deploying Google Kubernetes Engine
- Summary

Part 2 Kubernetes Architecture

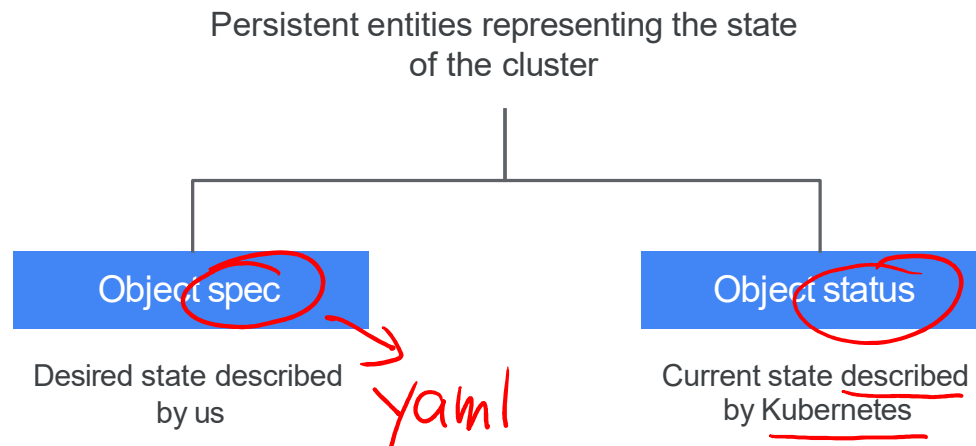
Kubernetes Concepts

Kubernetes objects

A Kubernetes object is defined as a persistent entity that represents the state of something running in a cluster: its desired state and its current state.

There are two elements to Kubernetes objects

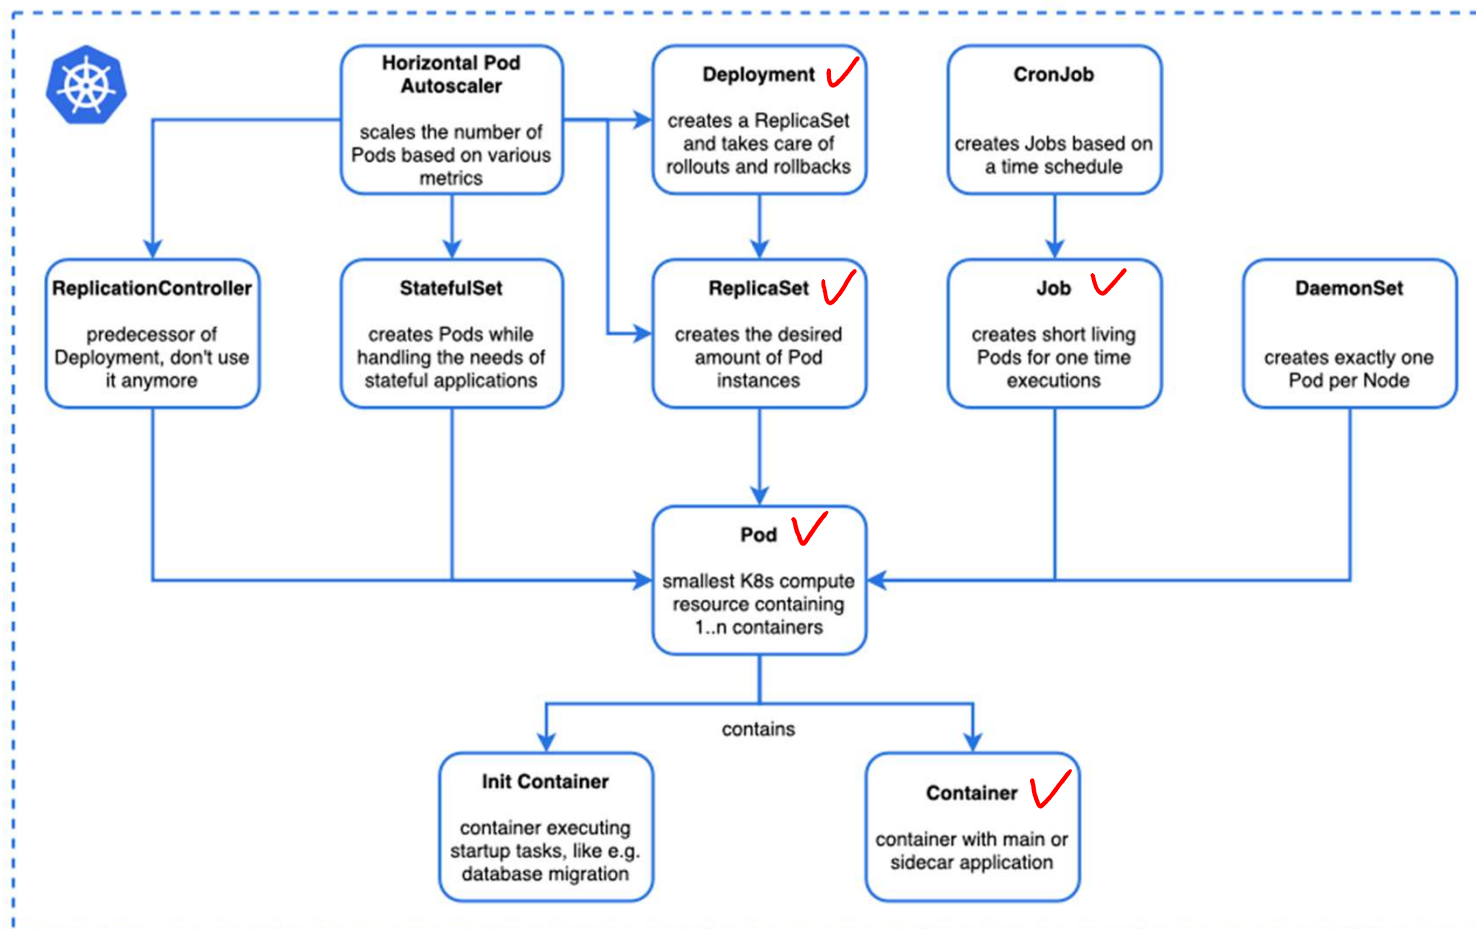
↑
yaml



Part 2 Kubernetes Architecture

Kubernetes Concepts

Kubernetes Workload Objects

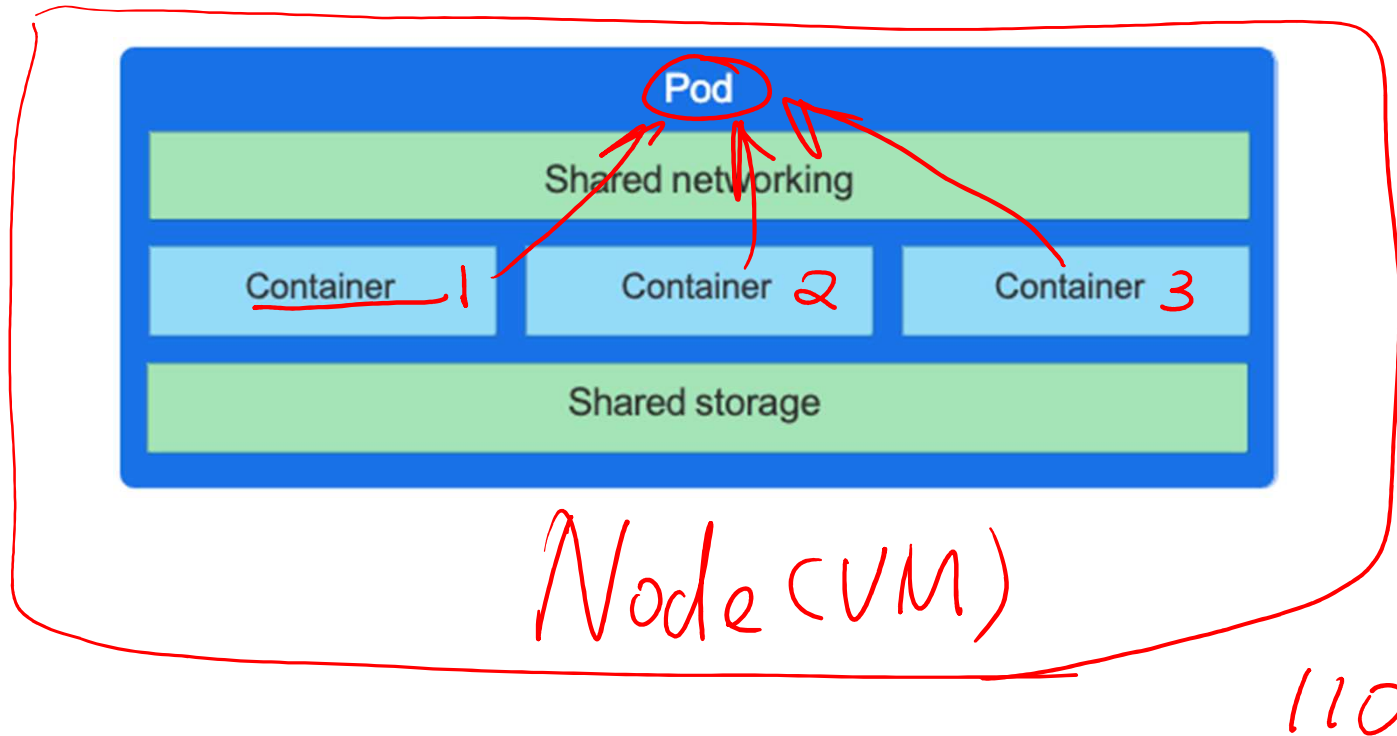


<https://shipit.dev/posts/kubernetes-overview-diagrams.html>

Part 2 Kubernetes Architecture

Kubernetes Concepts

Containers in a Pod share resources



Part 2 Kubernetes Architecture

Kubernetes Concepts

Running three nginx containers

You want three nginx containers running all the time

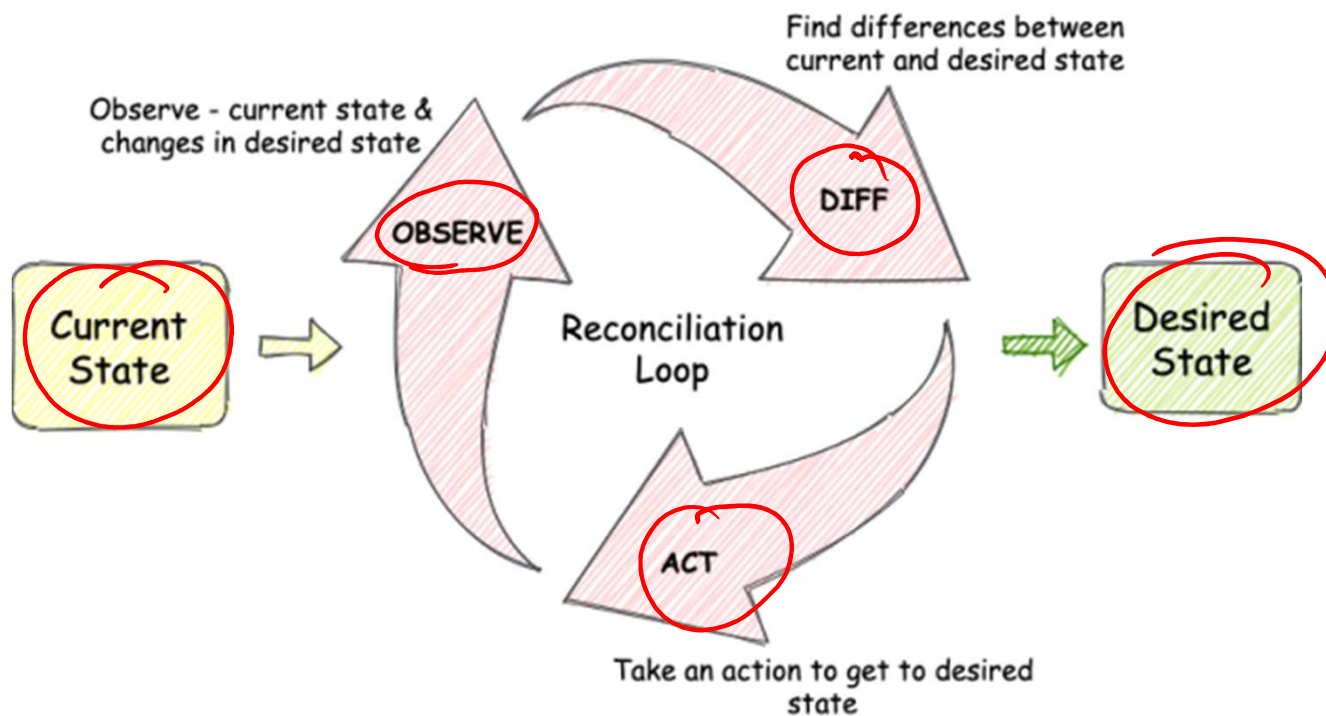
You declare objects that represent those containers

Kubernetes launches those objects and maintains them


Part 2 Kubernetes Architecture

Kubernetes Concepts

Desired state compared to current state

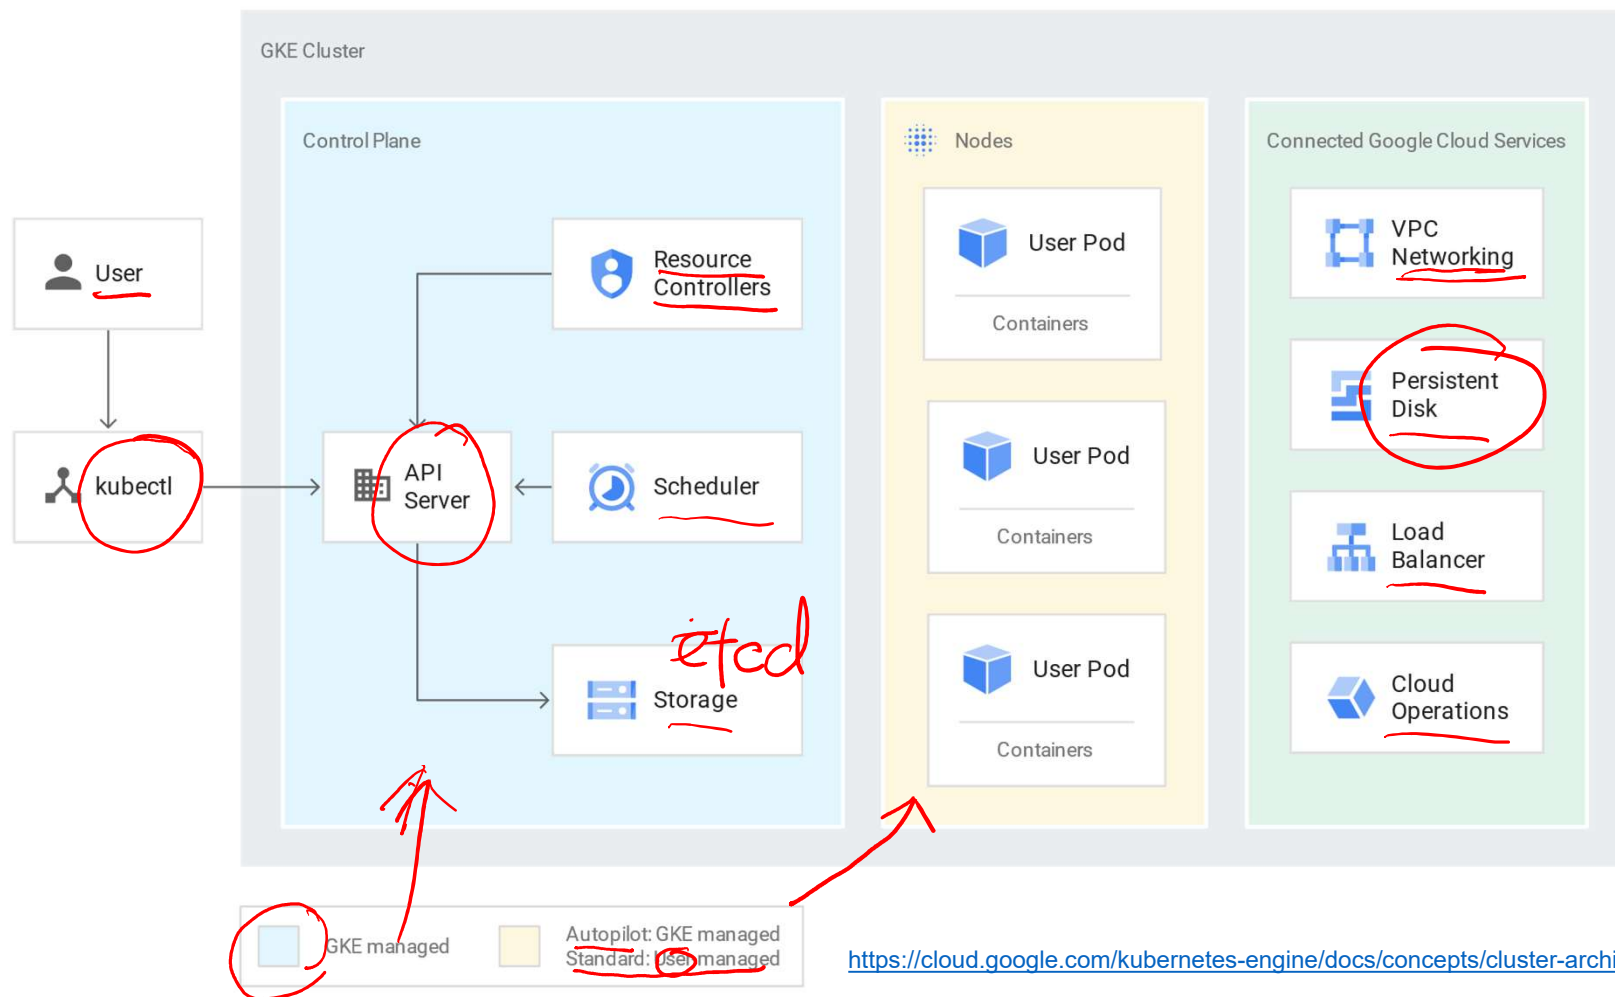


Part 2 Kubernetes Architecture



Part 2 Kubernetes Architecture

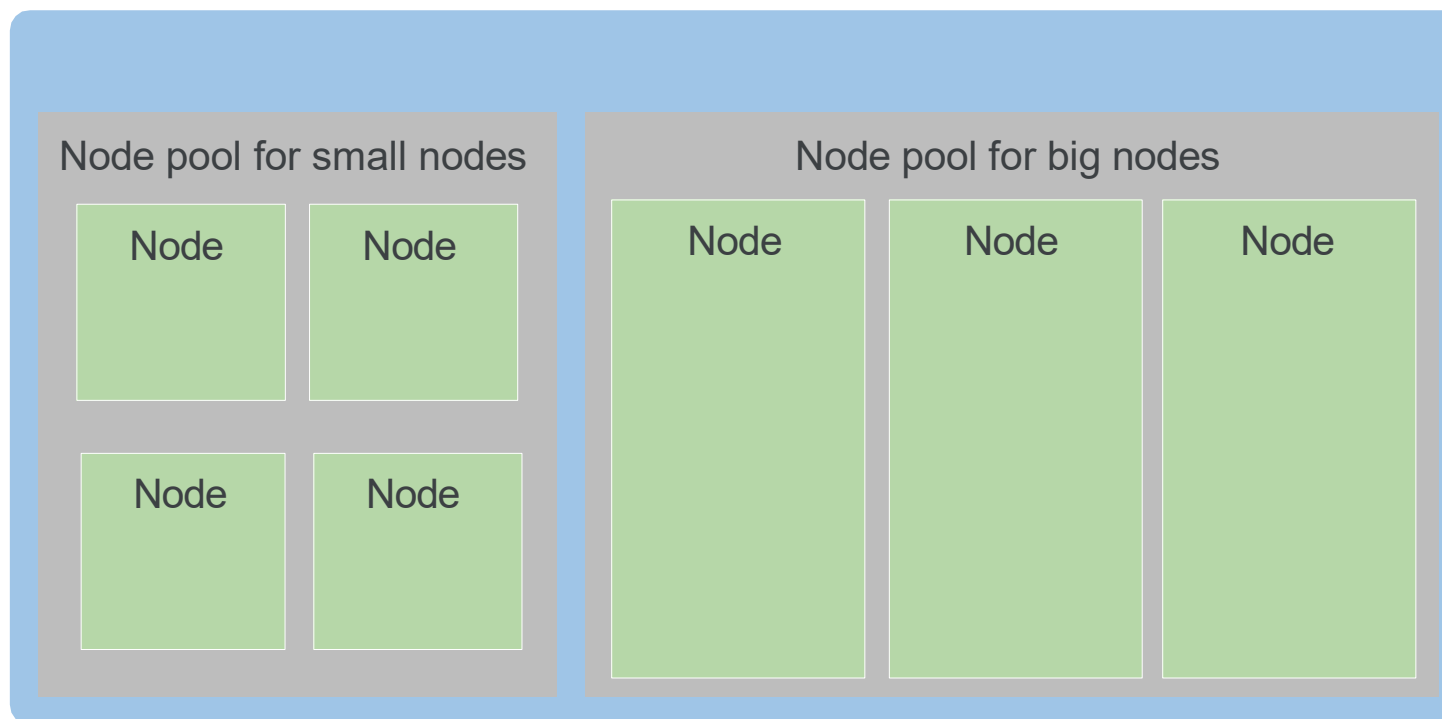
GKE Architecture



Part 2 Kubernetes Architecture

GKE Architecture

Node pools (GKE feature)

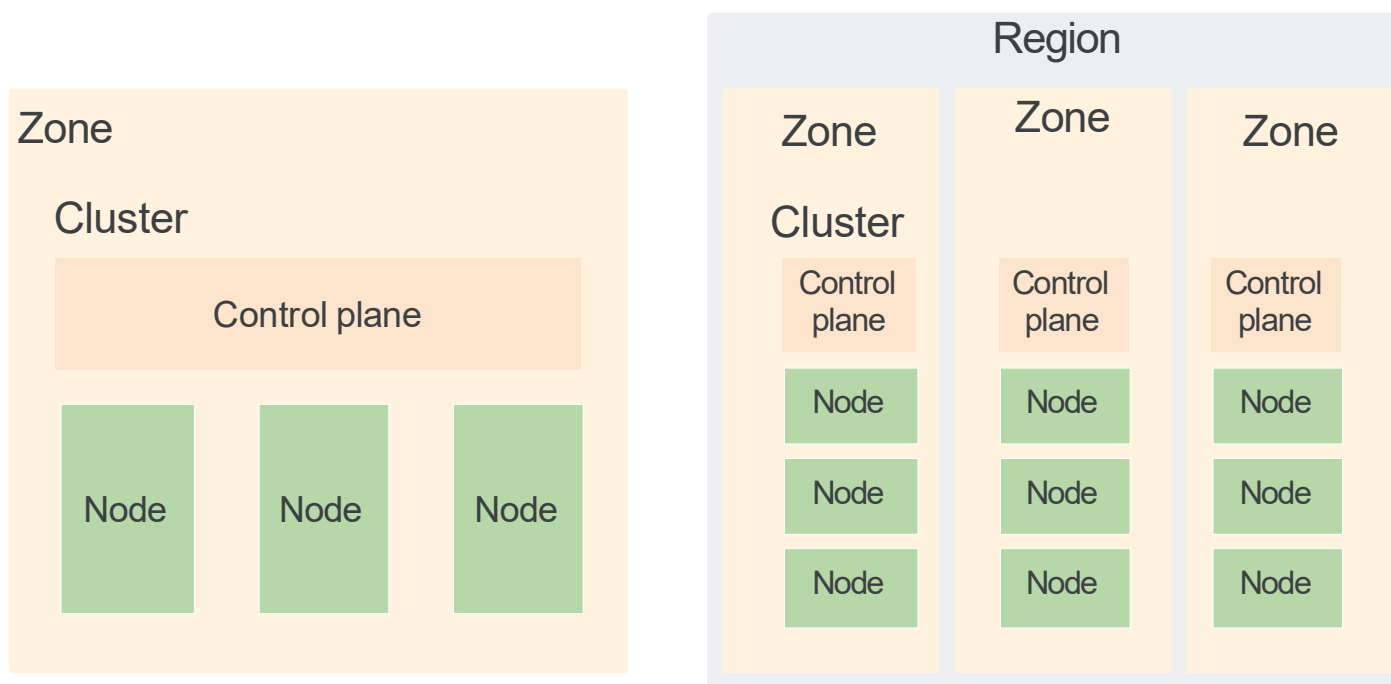


<https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-architecture>

Part 2 Kubernetes Architecture

GKE Architecture

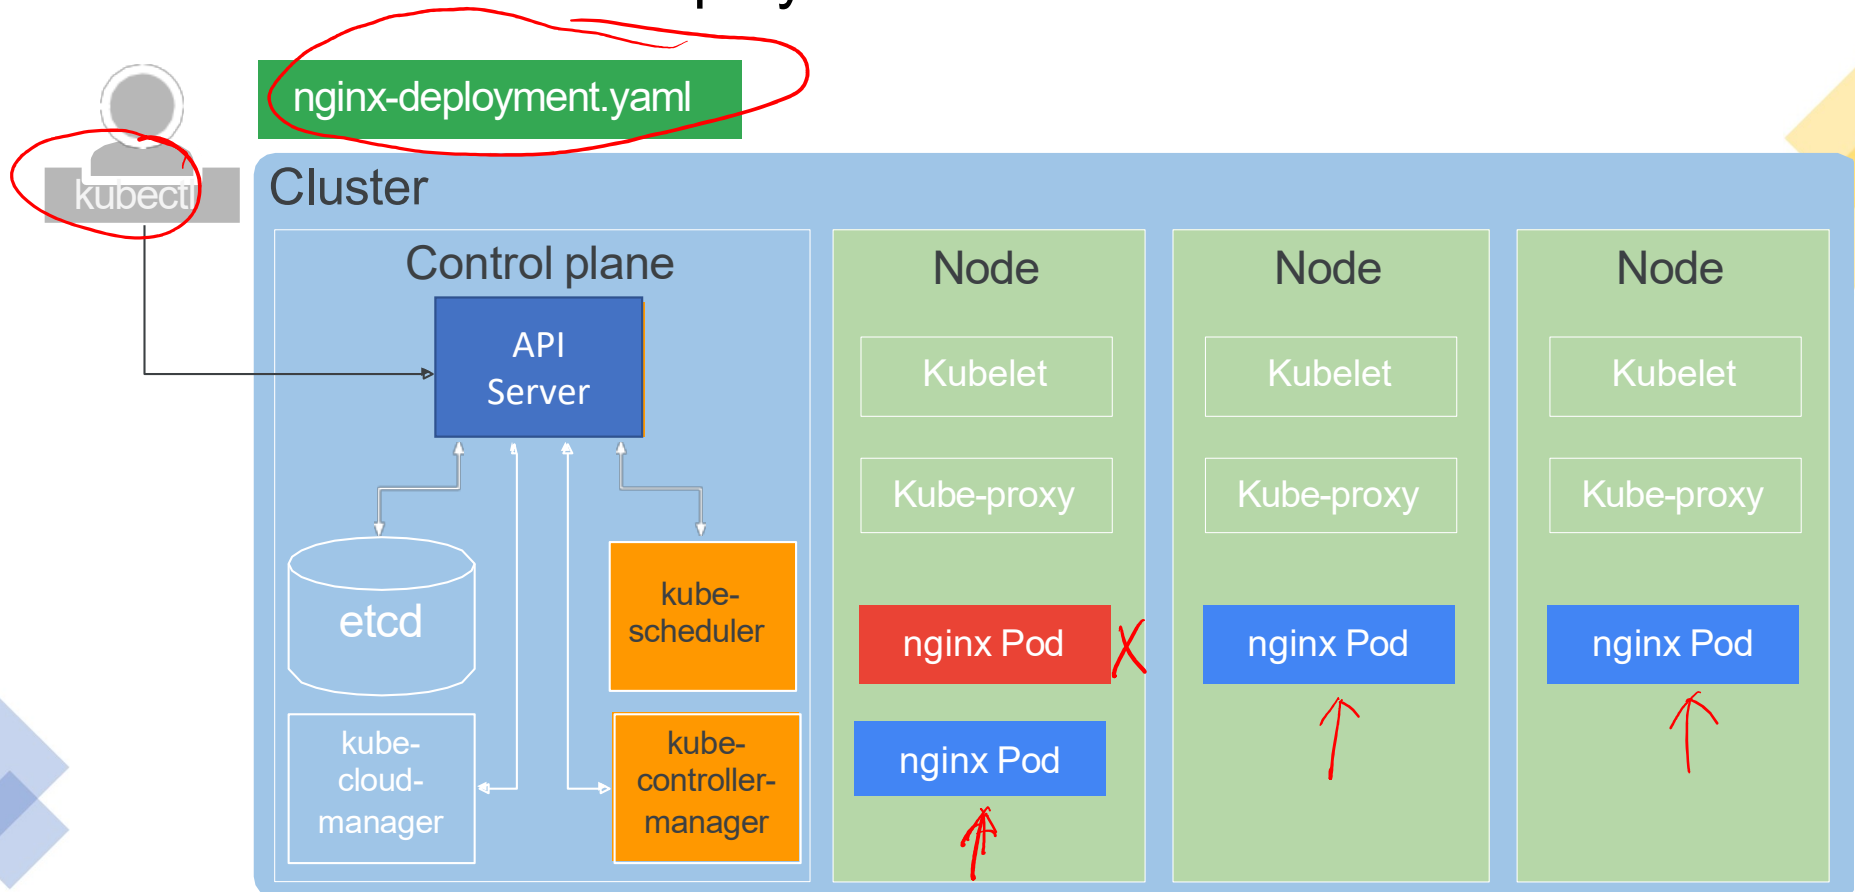
Zonal and regional clusters



Part 2 Kubernetes Architecture

GKE Architecture

Deployment of a workload

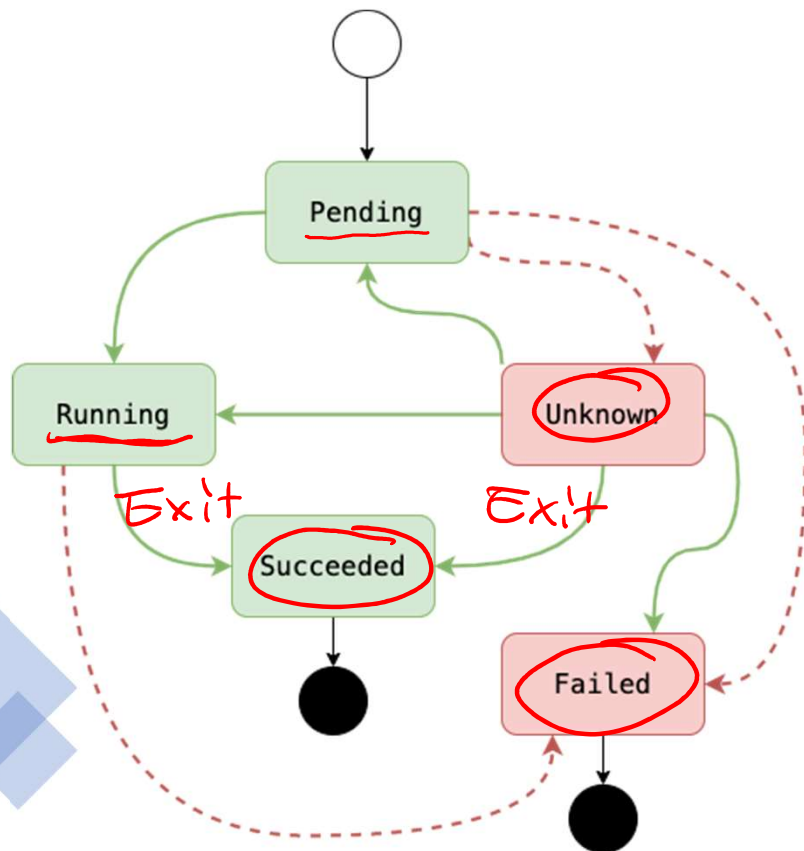


<https://kubernetes.io/docs/concepts/architecture/cloud-controller/>

Part 2 Kubernetes Architecture

GKE Architecture

The life cycle of a pod



Pod phases

Pending	Pod has been created by the cluster, but one or more of its containers are not yet running. This phase includes time spent being scheduled on a node and downloading images
Running	The Pod has been allotted to a node; all the containers have been created. At least one container is still running, or is in the process of starting or restarting
Succeeded	All containers in the Pod have terminated successfully
Failed	One or more containers terminated with non-zero status
Unknown	The state of the Pod cannot be determined. This occurs due to error while communicating with the node

Container states

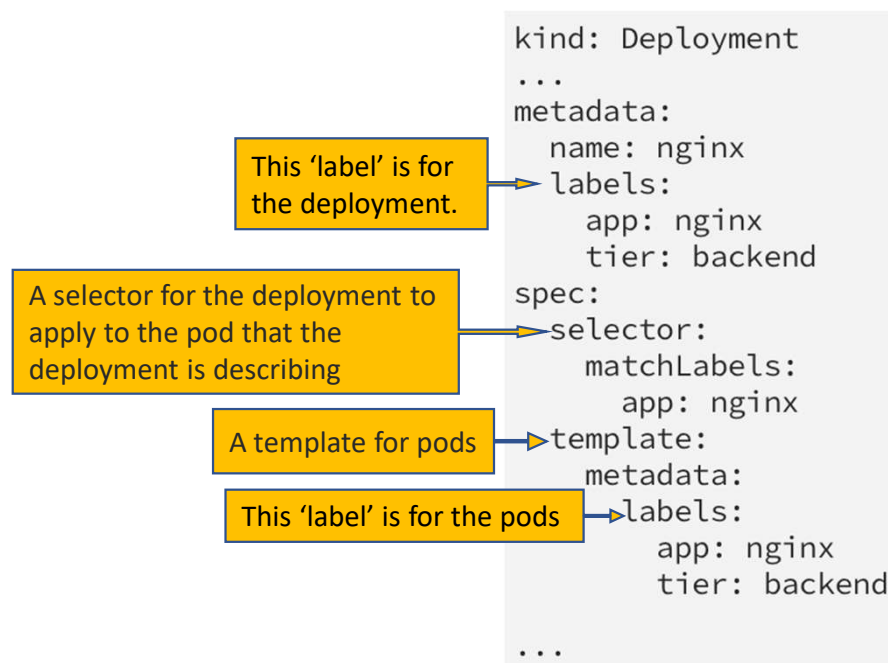
Waiting	When the container still <u>pulling image</u> , applying <u>Secret data etc.</u>
Running	When the container <u>executing without any issues</u>
Terminated	When the container <u>exited with non-zero status</u>

<http://millionvisit.blogspot.com/2021/03/kubernetes-for-developers-9-Kubernetes-Pod-Lifecycle.html>

Part 2 Kubernetes Architecture

GKE Architecture

Example yaml files



Part 2 Kubernetes Architecture

GKE Architecture


Create a GKE cluster and deploy a workload

<https://cloud.google.com/kubernetes-engine/docs/deploy-app-cluster>

<https://www.youtube.com/watch?v=p2LyoePiBo8>

doc





Part 2 Kubernetes Architecture

- Kubernetes Concepts
- Kubernetes Components
- Object Management
- Lab: Deploying Google Kubernetes Engine
- Summary

Part 2 Kubernetes Architecture

Object Management

Objects are defined in a YAML file
All objects are identified by a name

```
1  apiVersion: v1
2  kind: ReplicationController
3  metadata:
4    name: nginx
5  spec:
6    replicas: 3
7    selector:
8      app: nginx
9    template:
10     metadata:
11       name: nginx
12       labels:
13         app: nginx
14     spec:
15       containers:
16       - name: nginx
17         image: nginx
18         ports:
19         - containerPort: 80
20
```


```
hello.deployment.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: java-hello-world
5  spec:
6    replicas: 1
7    selector:
8      matchLabels:
9        app: java-hello-world
10   template:
11     metadata:
12       labels:
13         app: java-hello-world
14     spec:
15       containers:
16       - name: frontend
17         ima
```

image	string
imagePullPolicy	string
^↓ and ^↑ will move caret down and up in the editor Next Tip	

Part 2 Kubernetes Architecture

Object Management

Best practice tip: Use version control on YAML files



Part 2 Kubernetes Architecture

GKE Architecture

Example yaml file

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Kubernetes API version

Type of K8 Object i.e., Deployment

Name of the Deployment

No. of Pods to run at anytime

Type of Deployment Strategy

Maximum Pods to be created on top of desired count during deployment

Maximum Pods are unavailable during deployment

Both labels should match to create a Pod

Defining Pod manifest

Name of the container

Name of the container image

Application listening/running port

Deployment manifest: nginx-deployment.yaml


→ artifact registry

resources

<http://millionvisit.blogspot.com/2021/05/kubernetes-for-developers-14-Kubernetes-Deployment-YAML-manifest.html>

Resource management for Pods and Containers

- ✓ Important that containers have enough resources to run.
- ✓ Applications could use more resources than they should.
- ✓ CPU and memory (RAM) resources are the most common resources specified.



Part 2 Kubernetes Architecture

- Kubernetes Concepts
- Kubernetes Components
- Object Management
- Lab: Deploying Google Kubernetes Engine
- Summary

Part 2 Kubernetes Architecture

Summary

Summary

Kubernetes controllers keep the cluster state matching the desired state.

Kubernetes consists of a family of control plane components, running on the control plane and the nodes.

GKE abstracts away the control plane.

Declare the state you want using manifest files.

yaml