




CSCI 5409 Cloud Computing  
Fall, 2023  
Instructor: Dr. Lu Yang

Kubernetes Workload &  
Production (2)  
Oct 16, 2023

## Housekeeping and Feedback

- Start recording
- Midterm in class next Monday, Oct 23.
  - 10 multiple choice, 4 short answer, and 2 long answer
  - It covers the contents up to this lecture.




## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Lab

Creating Google Kubernetes Engine Deployments  
(<https://www.youtube.com/watch?v=k4x4ce370LA>)



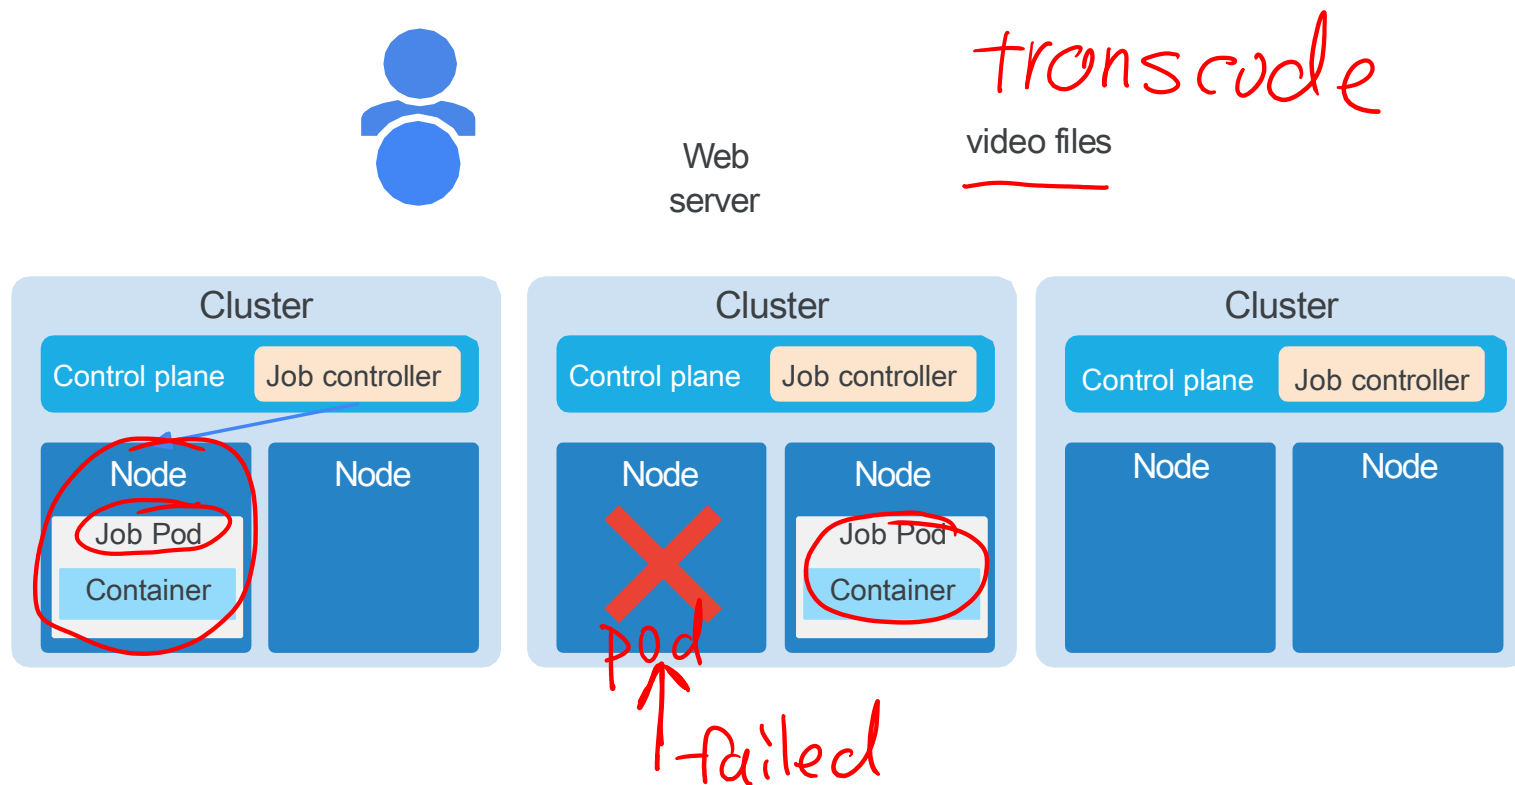
## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- **Jobs**
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Jobs

A scenario where Job provides the solution



## Part 2 Deployments, jobs, and scaling

### Jobs

### Jobs definition

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-app-job
spec:
  completions: 3
  template:
    spec:
[ ... ]
```

## Part 2 Deployments, jobs, and scaling

### Jobs

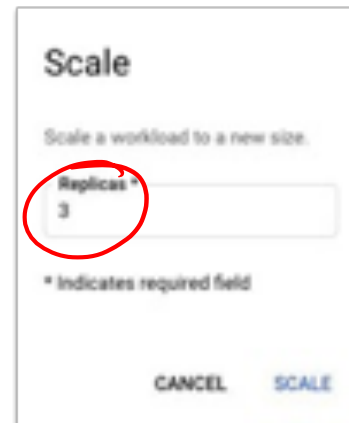
#### Inspecting a job

```
$ kubectl describe job [JOB_NAME]
```

```
$ kubectl get pod -l [job-name=my-app-job]
```

#### Scaling a job

```
$ kubectl scale job [JOB_NAME] --replicas [VALUE]
```



The screenshot shows a 'Scale' dialog box with the title 'Scale'. Below the title is the instruction 'Scale a workload to a new size.' There is a text input field labeled 'Replicas \*' with the value '3' entered. A red circle highlights the 'Replicas \*' label and the input field. Below the input field is the text '\* Indicates required field'. At the bottom of the dialog are two buttons: 'CANCEL' and 'SCALE'.



## Part 2 Deployments, jobs, and scaling

### Jobs

#### Deleting a job

```
$ kubectl delete -f [JOB_FILE]
```

```
$ kubectl delete job [JOB_NAME]
```

## Part 2 Deployments, jobs, and scaling

### Jobs


### Differences of jobs and deployments

The main difference between Deployments and Jobs is **how they handle a Pod that is terminated**. A Deployment is intended to be a "service", e.g. it should be up-and-running, so it will try to restart the Pods it manage, to match the desired number of replicas. While a Job is intended to execute and successfully terminate.

e.g. webserver,  
database server

e.g. database  
backup

In a Deployment, the default restartPolicy of your Pod is set Always. In a Job: Never. A job is not meant to restart your container once it would have completed. A deployment is not meant to complete.



## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Lab

Deploying jobs on GKE  
(<https://www.youtube.com/watch?v=pqGfzXHrYLk>)



## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools

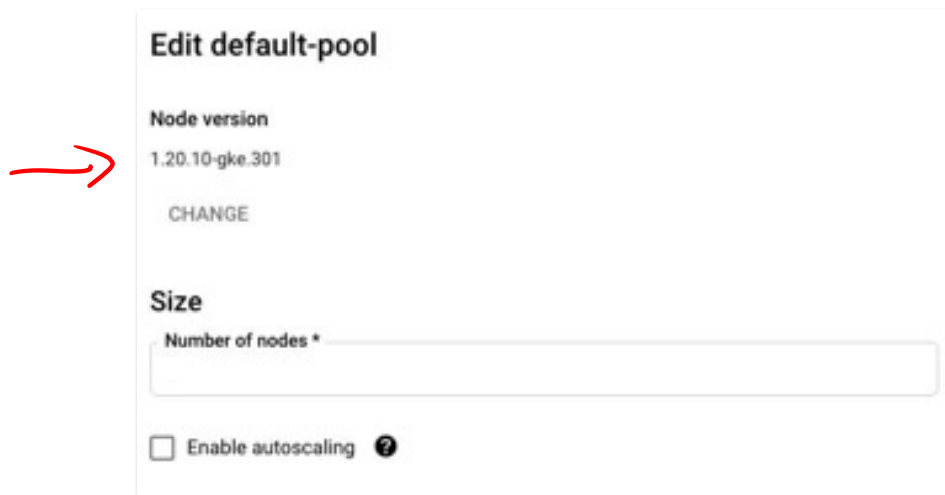
## Part 2 Deployments, jobs, and scaling

### Cluster scaling

Scaling down a cluster using the gcloud command

```
gcloud container clusters resize projectdemo --  
node-pool \ default-pool \  
--num_nodes 6
```

Scaling down a cluster from the cloud console



**Edit default-pool**

**Node version**  
1.20.10-gke.301  
CHANGE

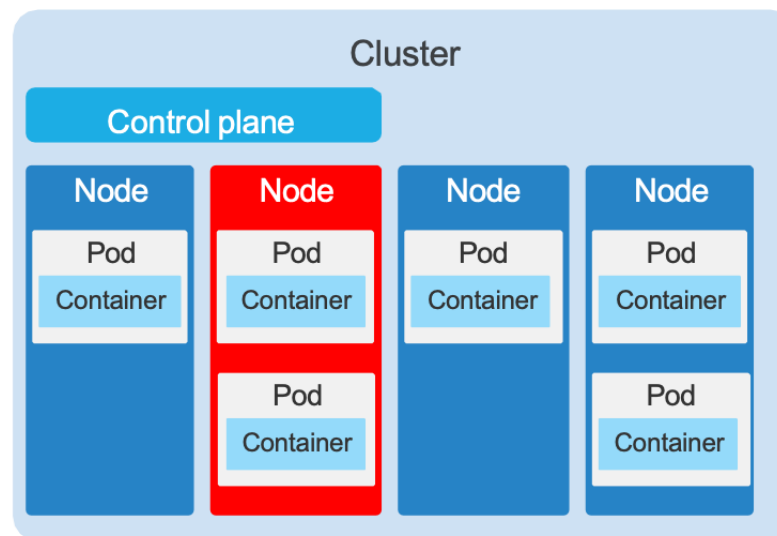
**Size**  
Number of nodes \*

☐ Enable autoscaling ?

## Part 2 Deployments, jobs, and scaling

### Cluster scaling

Manual Cluster Scale down selects nodes randomly

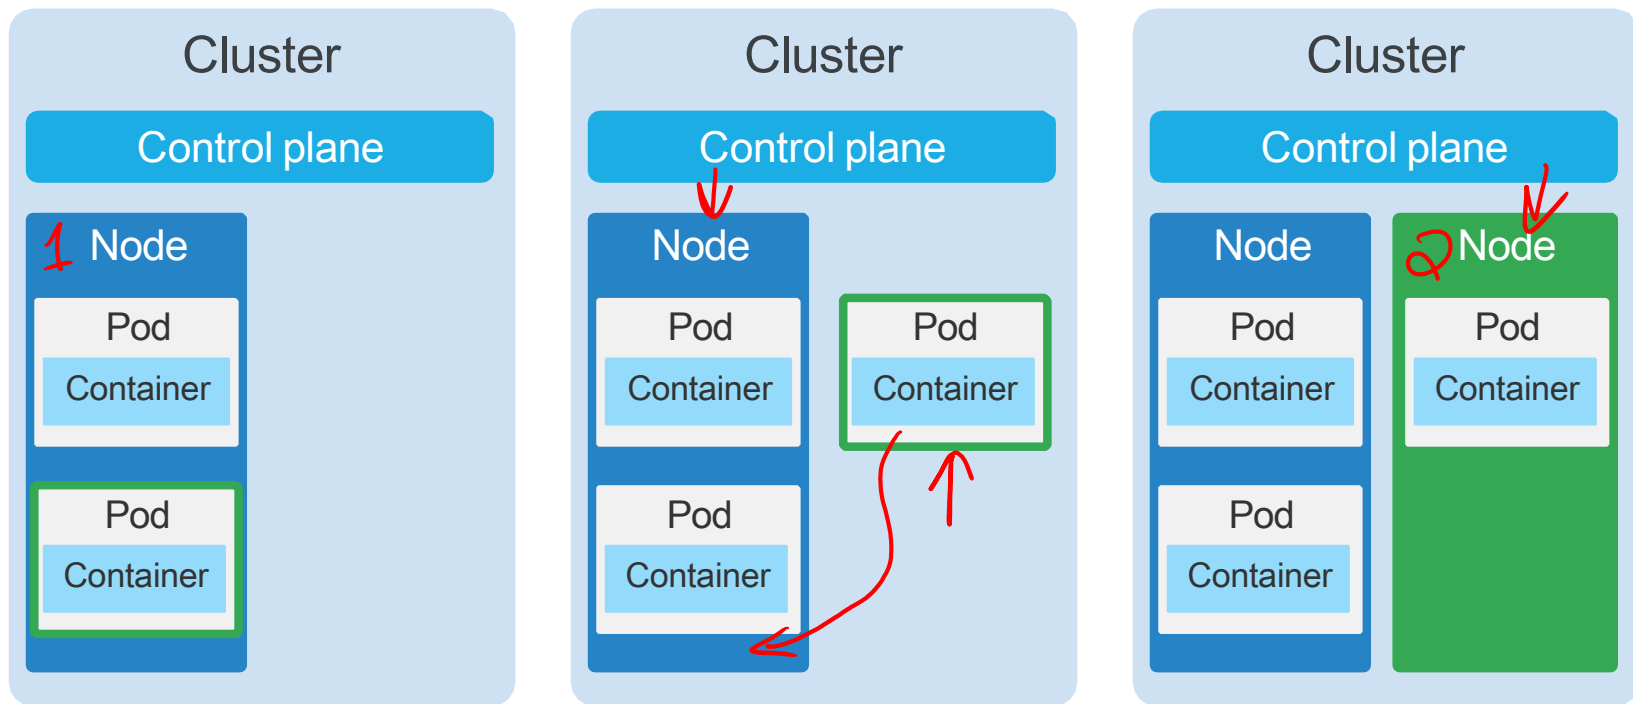


## Part 2 Deployments, jobs, and scaling

### Cluster scaling

Scale up a cluster with autoscaling

*ECS ↔ EKS*

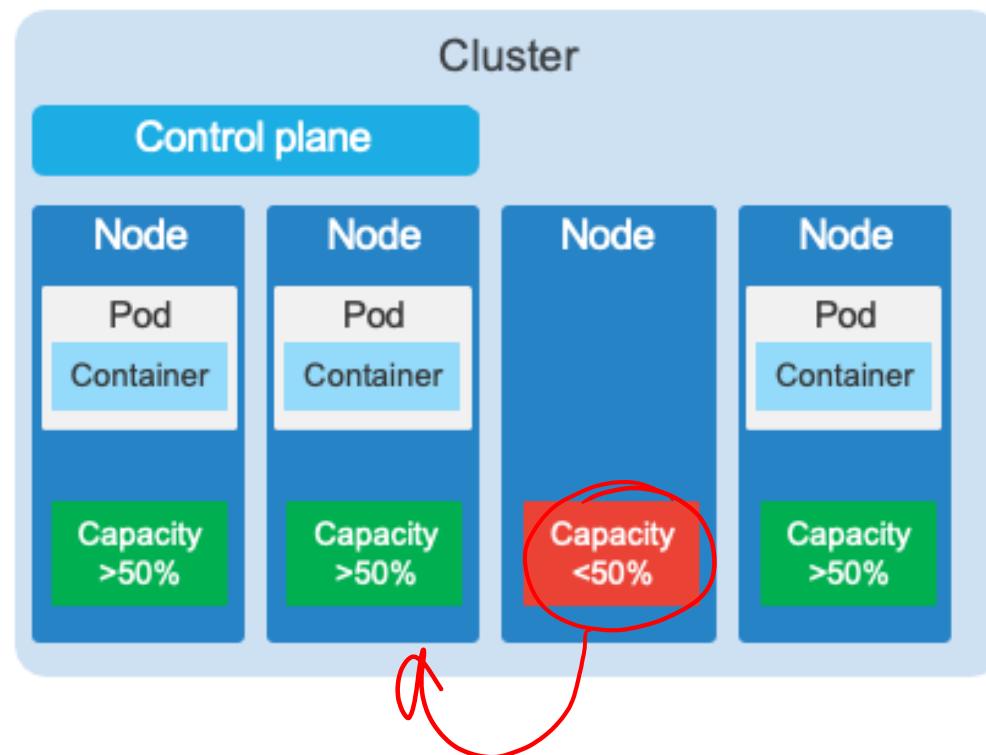




## Part 2 Deployments, jobs, and scaling

### Cluster scaling

The Autoscaler removes nodes that remain below 50% utilization



## Part 2 Deployments, jobs, and scaling

### Cluster scaling

GPU  
pool

Setting a node pool size




Node pool = 0

Cluster size  $\neq 0$

MAX = 15,000 nodes x 110 Pods

Increase quota limits to avoid disruption

High Mem



## Part 2 Deployments, jobs, and scaling

### Cluster scaling

### gcloud commands for autoscaling

Create a cluster with  
autoscaling enabled

```
gcloud container clusters create  
[CLUSTER_NAME] --enable-autoscaling  
--min-nodes 15 --max-nodes 50  
[--zone COMPUTE_ZONE]
```

Enable autoscaling for an  
existing node pool


```
gcloud container clusters update  
[CLUSTER_NAME] --enable-autoscaling \  
--min-nodes 1 --max-nodes 10 --zone  
[COMPUTE_ZONE] --node-pool [POOL_NAME]
```

Add a node pool with  
autoscaling enabled

```
gcloud container node-pools create  
[POOL_NAME] --cluster [CLUSTER_NAME]  
--enable-autoscaling --min-nodes 15  
--max-nodes 50 [--zone COMPUTE_ZONE]
```

Disable autoscaling for an  
existing node pool

```
gcloud container clusters update  
[CLUSTER_NAME] --no-enable-autoscaling \  
--node-pool [POOL_NAME] [--zone  
[COMPUTE_ZONE] --project [PROJECT_ID]]
```




## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools

## Part 2 Deployments, jobs, and scaling

### Controlling pod placement

### Controlled scheduling



## Part 2 Deployments, jobs, and scaling

### Controlling pod placement

Nodes must match all the labels present  
under the nodeSelector field

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql
  labels:
    env: test

  containers:
  - name: mysql
    image: mysql
    imagePullPolicy: IfNotPresent
  nodeSelector:
    disktype: ssd
[...]
```

```
apiVersion: v1
kind: Node
metadata:
  name: node1
  labels:
    disktype: ssd
[...]
```

## Part 2 Deployments, jobs, and scaling

### Controlling pod placement

Nodes must match all the labels present  
under the nodeSelector field

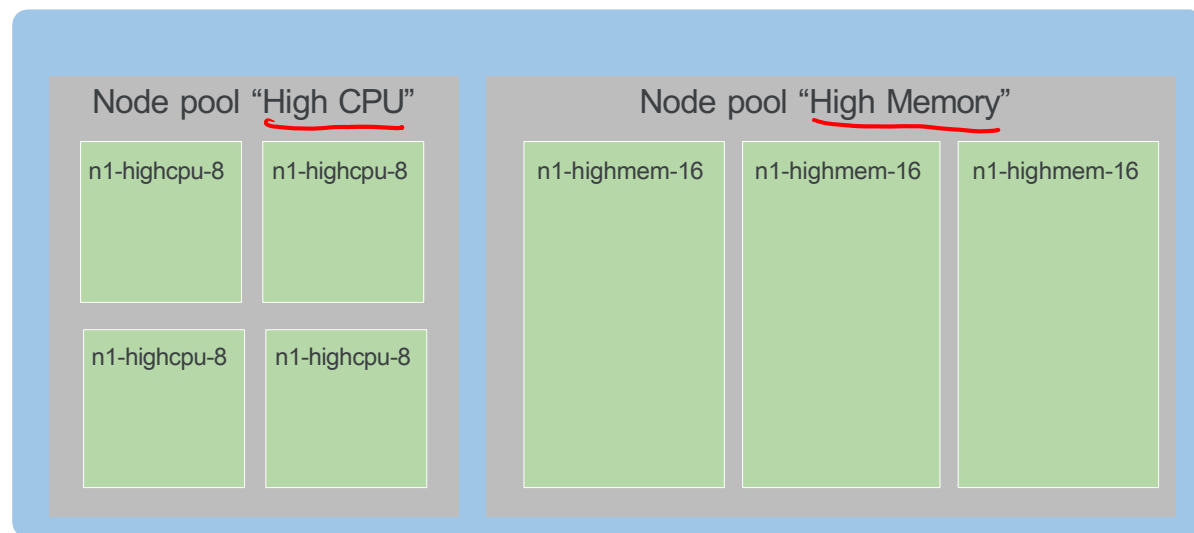
*auto generated  
labels*

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql
  labels:
    env: test
spec:
  containers:
    - name: mysql
      image: mysql
      imagePullPolicy: IfNotPresent
  nodeSelector:
    cloud.google.com/gke-nodepool=ssd
[...]
```


## Part 2 Deployments, jobs, and scaling

### Controlling pod placement

Use node pools to manage different kinds of nodes







## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Getting software into your cluster


#### How to get software

- Build it yourself, and supply your own YAML.
- Use Helm to install software into your cluster.
- Use Google Cloud Marketplace to install both open-source and commercial software.

## Part 2 Deployments, jobs, and scaling

Getting software into your cluster

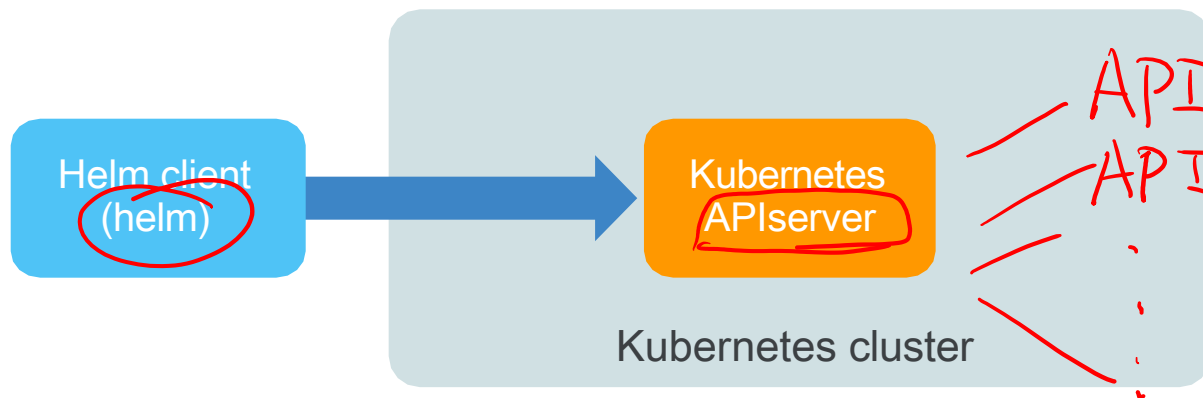
Organize Kubernetes objects in packages and deploy complex packages



## Part 2 Deployments, jobs, and scaling

Getting software into your cluster

Helm interacts directly with the Kubernetes APIserver






## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Lab

Configuring Pod Autoscaling and Node Pools  
(<https://www.youtube.com/watch?v=TDuBmjZqpPQ>)



## Part 2 Deployments, jobs, and scaling

- Deployments
- Self-learning lab: Creating Google Kubernetes Engine Deployments
- Jobs
- Self-learning lab: Deploying Jobs on Google Kubernetes Engine
- Cluster Scaling
- Controlling Pod Placement
- Getting Software into Your Cluster
- Self-learning lab: Configuring Pod Autoscaling and Node Pools
- Summary

## Part 2 Deployments, jobs, and scaling

### Summary

#### Summary


- Create and use Deployments.
- Create and run Jobs.
- Use Helm Charts.
- Scale clusters manually and automatically.





## Part 3 Persistent data and storage

- Volumes
- Self-learning lab: Configuring Persistent Storage for Google Kubernetes Engine
- Summary



## Part 3 Persistent data and storage

- Volumes
- Self-learning lab: Configuring Persistent Storage for Google Kubernetes Engine
- Summary

## Part 3 Persistent data and storage

### Volumes

Kubernetes offers storage abstraction options

#### Volumes

Are a directory which is accessible to all of the containers in a Pod.

Some Volumes are ephemeral. ↙

Some Volumes are persistent.

#### Persistent Volumes

Manage durable storage in a cluster.


Are independent of the Pod's lifecycle.

Provisioned dynamically through PersistentVolumeClaims or explicitly created by a cluster admin.

## Part 3 Persistent data and storage

### Volumes

Volumes allow containers within a Pod to share data



Short-lived



Long-lived

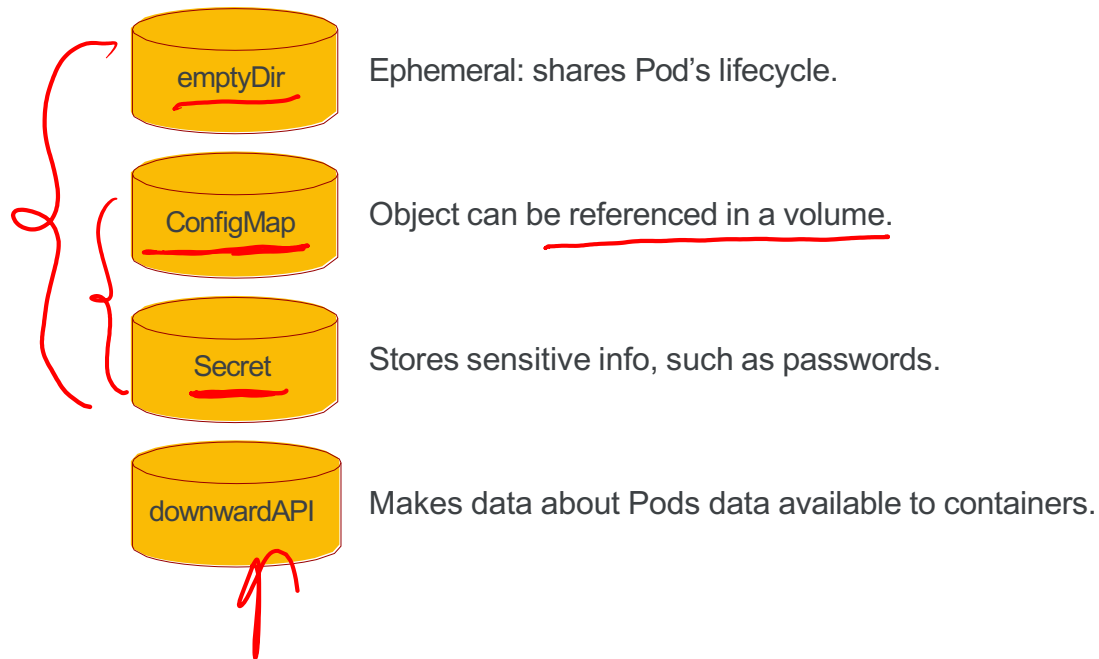


Volume specifications

## Part 3 Persistent data and storage

### Volumes


#### Ephemeral volume types



## Part 3 Persistent data and storage

### Volumes

emptyDir volume:  
created when a Pod is assigned to a node

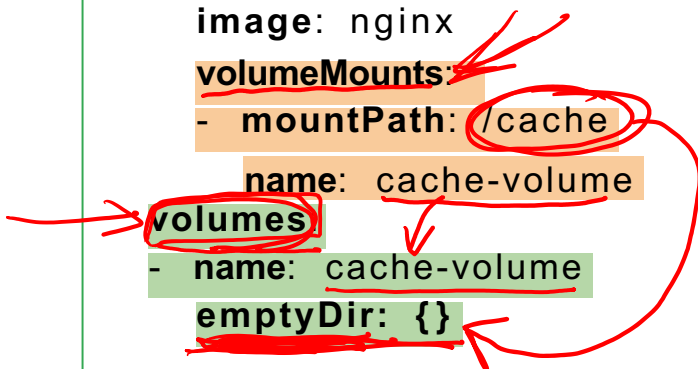


## Part 3 Persistent data and storage

### Volumes

#### Creating a Pod with an emptyDir volume

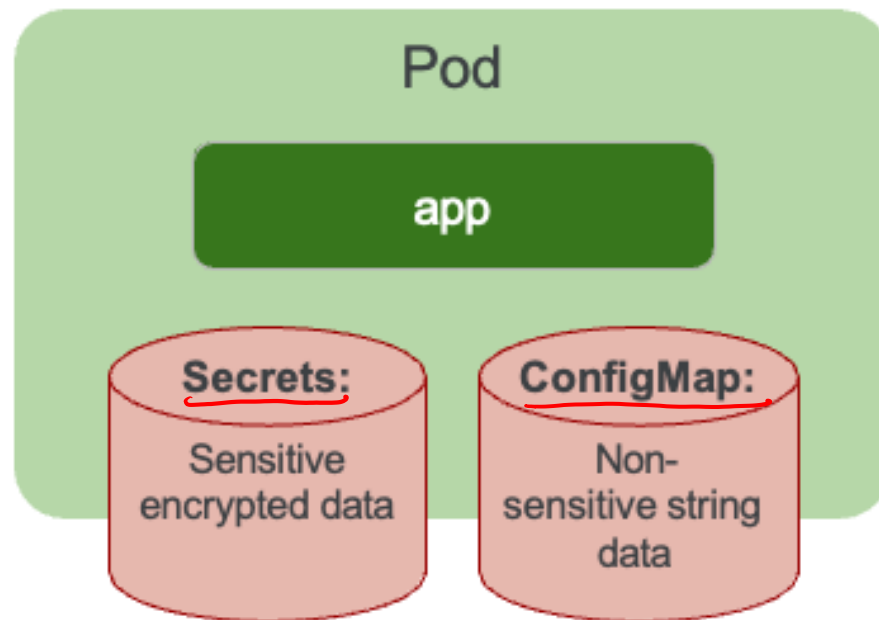
```
apiVersion: v1
kind: Pod
metadata:
  name: web
spec:
  containers:
    - name: web
      image: nginx
      volumeMounts:
        - mountPath: /cache
          name: cache-volume
  volumes:
    - name: cache-volume
      emptyDir: {}
```



## Part 3 Persistent data and storage

### Volumes

Secret and ConfigMap Volumes are ephemeral



Lab: [https://www.youtube.com/watch?v=BhC\\_-qqRdRk](https://www.youtube.com/watch?v=BhC_-qqRdRk)



## Part 4 Persistent data and storage

### Volumes

#### The benefits of PersistentVolumes

PD

PV

Abstracts storage  
provisioning from storage  
consumption.

Promotes microservices  
architecture.

Allows cluster administrators  
to provision and maintain  
storage.

Developers can claim  
provisioned storage for app  
consumption.

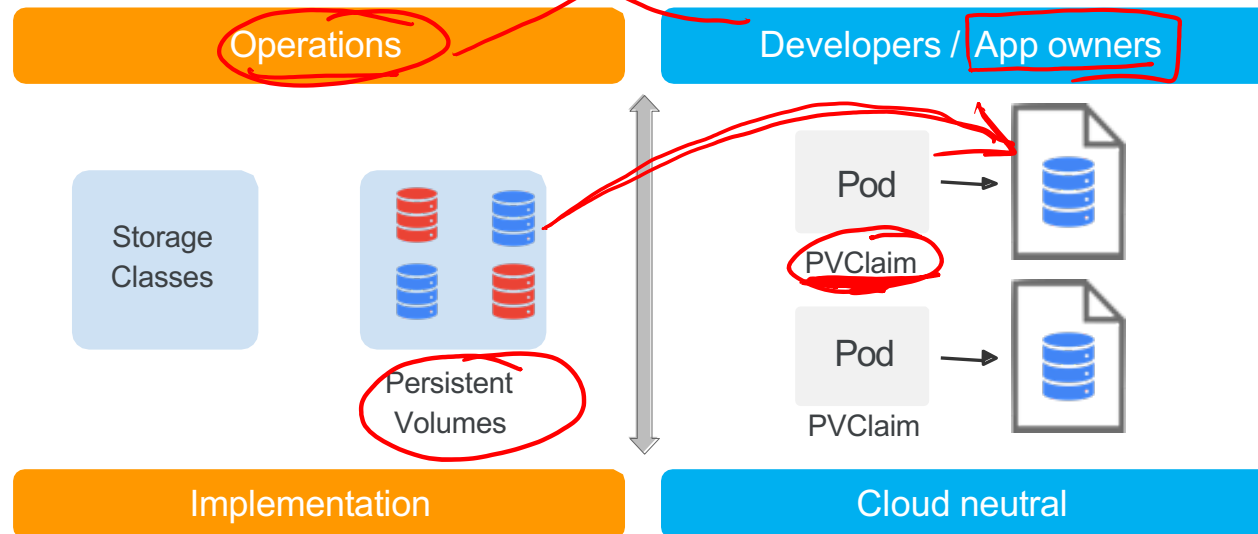
## Part 3 Persistent data and storage

### Volumes

*Consumer*

*Admin*

PersistentVolumeClaims and PersistentVolumes  
separate storage consumption from provisioning



## Part 3 Persistent data and storage

### Volumes

Creating a Compute Engine persistent disk using a gcloud command

PD

```
$ gcloud compute disks create  
--size=100GB  
--zone=us-central1-a demo-disk
```

## Part 3 Persistent data and storage


### Volumes

PersistentVolumes abstraction has two components

PersistentVolume (PV)

- Independent of a Pod's lifecycle.
- Managed by Kubernetes.
- Manually or dynamically provisioned.
- Persistent Disks are used by GKE as PersistentVolumes.

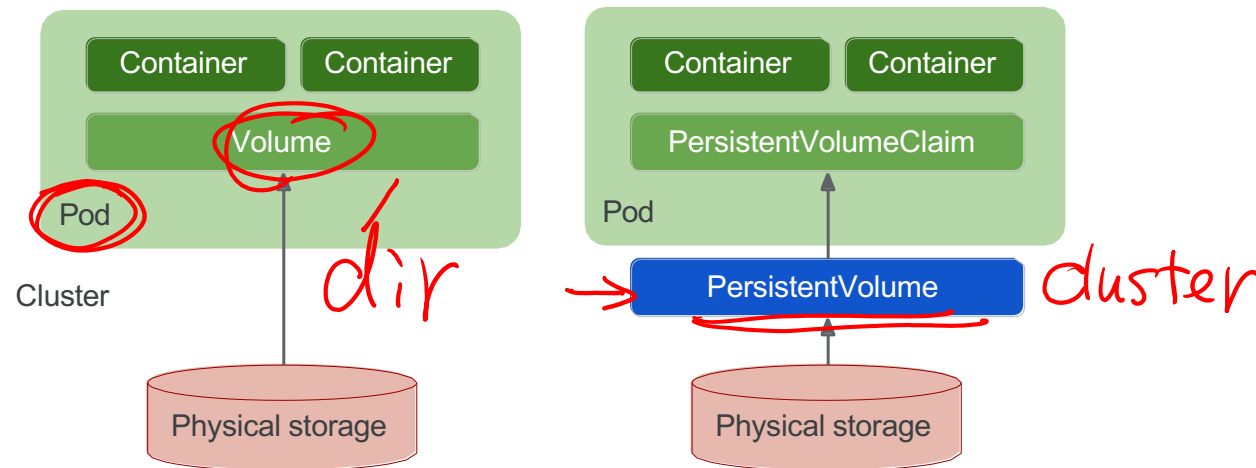
PersistentVolumeClaim (PVC)



## Part 3 Persistent data and storage

### Volumes

PersistentVolumes must be claimed



## Part 3 Persistent data and storage

### Volumes

The AccessModes you specify determine how this Volume can be read from or written to

<pre>apiVersion: v1 kind: PersistentVolume metadata:   name: pd-volume spec:   storageClassName: "standard"   capacity:     storage: 100G   accessModes:     - ReadWriteOnce   gcePersistentDisk:     pdName: demo-disk     fsType: ext4</pre>	<pre>apiVersion: v1 kind: PersistentVolume metadata:   name: pd-volume spec:   storageClassName: "standard"   capacity:     storage: 100G   accessModes:     - ReadOnlyMany   gcePersistentDisk:     pdName: demo-disk     fsType: ext4</pre>	<pre>apiVersion: v1 kind: PersistentVolume metadata:   name: pd-volume spec:   storageClassName: "nfs"   capacity:     storage: 100G   accessModes:     - ReadWriteMany   nfs:     path: /tmp     server: 172.17.0.2</pre>
--	---	--

node

- node

- node

cluster admin

PV & GKE  
AZ

## Part 3 Persistent data and storage

### Volumes

You can create a Persistent Volume from a YAML manifest

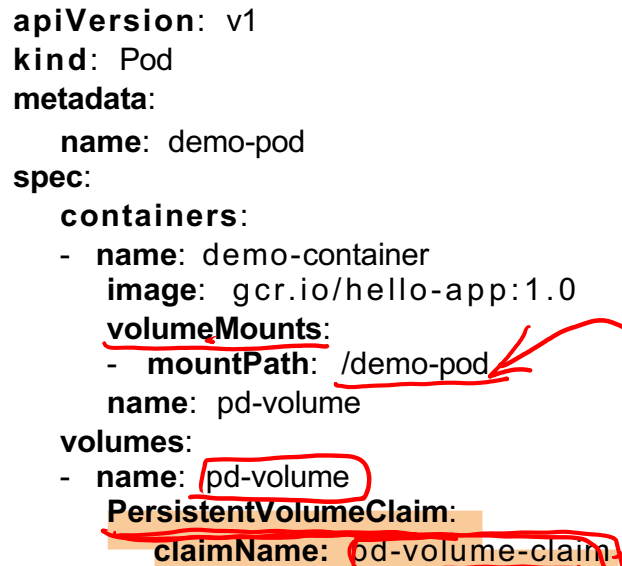
```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pd-volume
spec:
  storageClassName: "standard"
  capacity:
    storage: 100G
  accessModes:
    - ReadWriteOnce:
  gcePersistentDisk:
    pdName: demo-disk
    fsType: ext4
```

## Part 3 Persistent data and storage

### Volumes

The modern, easier-to-manage way is to use the PersistentVolume abstraction

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-pod
spec:
  containers:
    - name: demo-container
      image: gcr.io/hello-app:1.0
      volumeMounts:
        - mountPath: /demo-pod
          name: pd-volume
  volumes:
    - name: pd-volume
      PersistentVolumeClaim:
        claimName: pd-volume-claim
```






## Part 3 Persistent data and storage

### Volumes

The PersistentVolume can be retained when the PersistentVolumeClaim is deleted

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pd-volume-claim
spec:
  storageClassName: "standard"
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100G
  persistentVolumeReclaimPolicy: Retain
```



## Part 3 Persistent data and storage


- Volumes
- Self-learning lab: Configuring Persistent Storage for Google Kubernetes Engine
- Summary

## Part 3 Persistent data and storage

Lab

Configuring Persistent Storage  
for Google Kubernetes Engine

([https://www.youtube.com/watch?v=MaN\\_deRwrhs](https://www.youtube.com/watch?v=MaN_deRwrhs) )



## Part 3 Persistent data and storage

- Volumes
- Self-learning lab: Configuring Persistent Storage for Google Kubernetes Engine
- Summary

## Part 3 Persistent data and storage

### Summary

#### Summary

Understand and work with Kubernetes storage abstractions.

Use ConfigMaps to decouple configuration from Pods.

Manage and store sensitive authorization and authentication data.