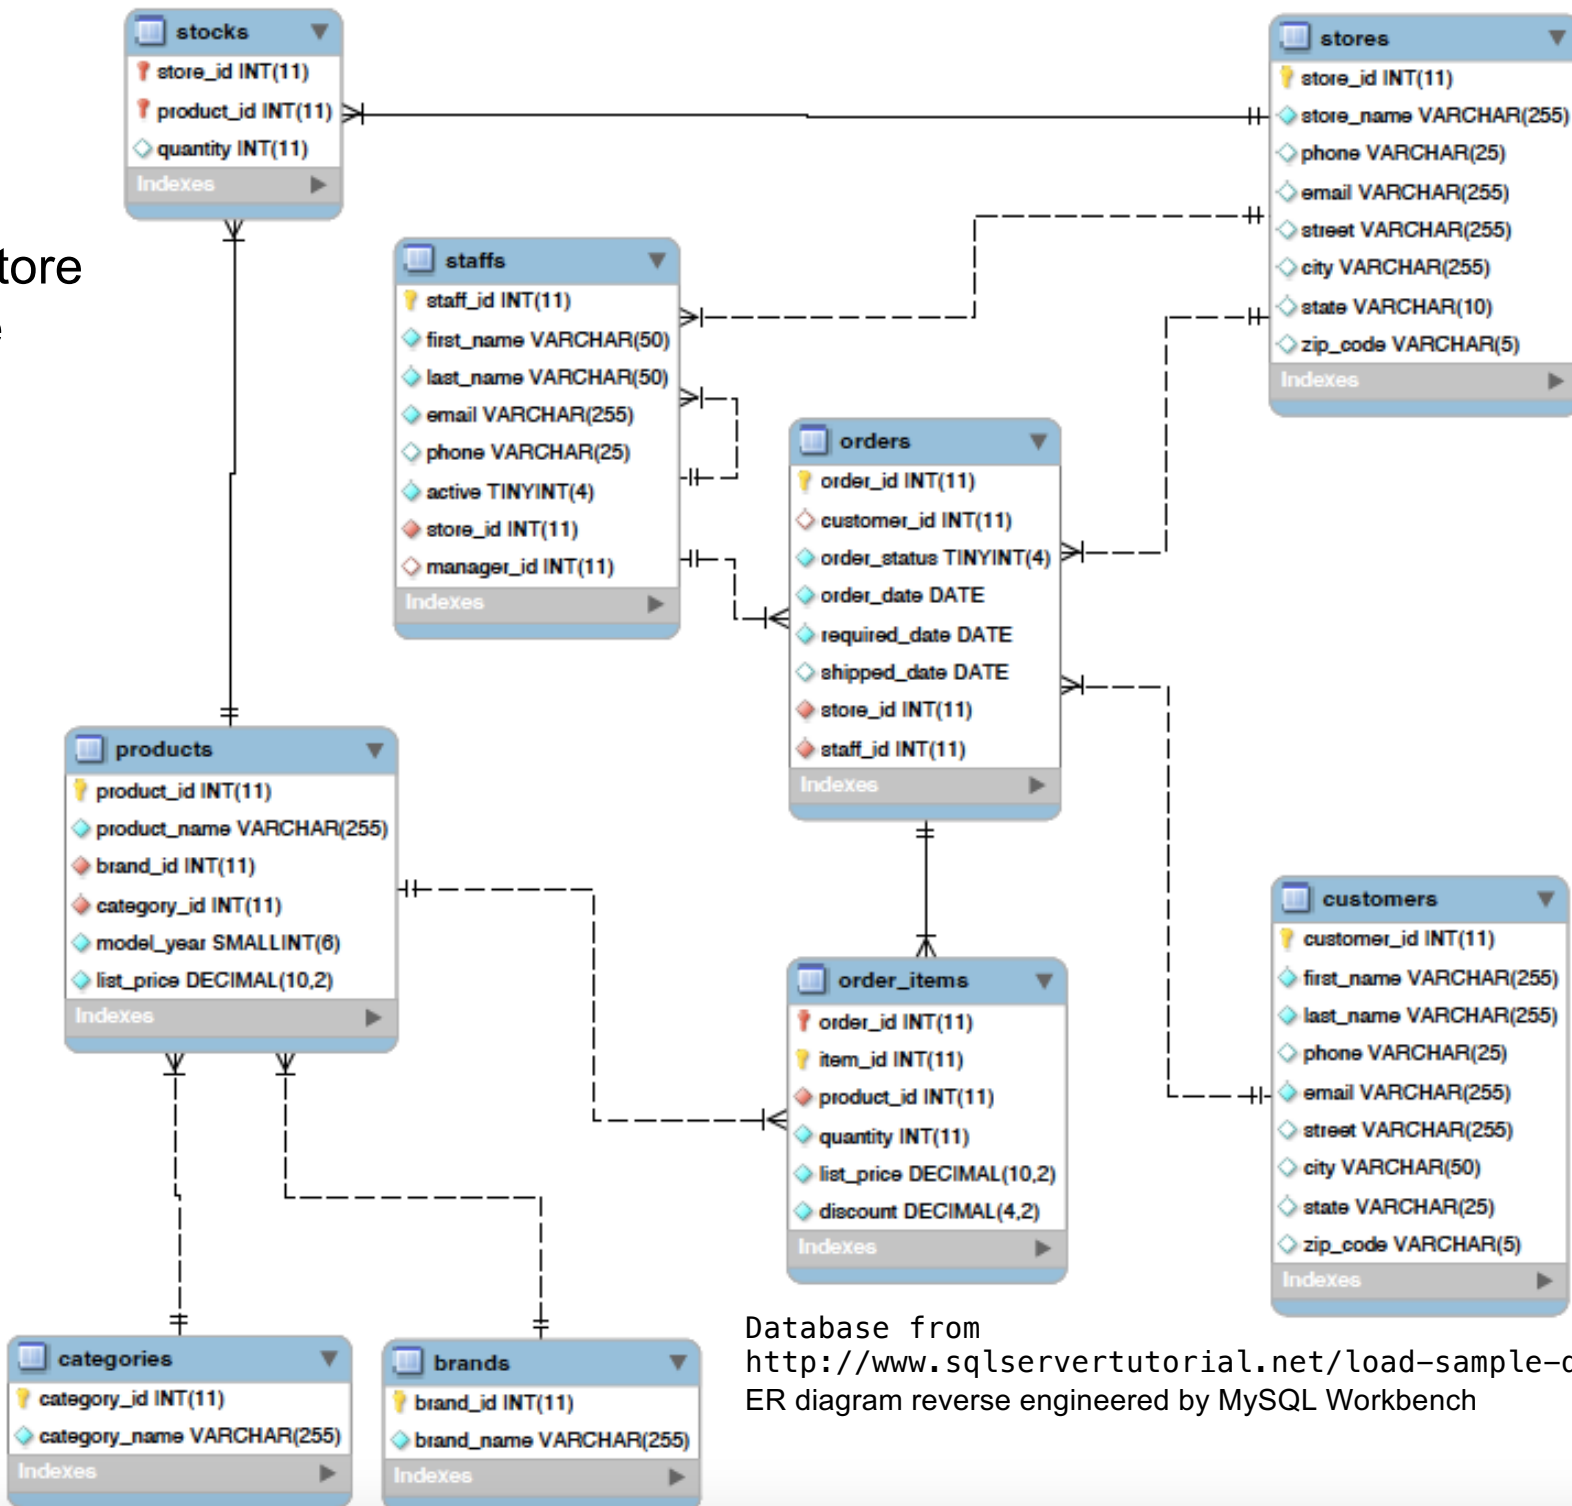


Bicycle store database



Database from
<http://www.sqlservertutorial.net/load-sample-database/>
 ER diagram reverse engineered by MySQL Workbench

Questions to try...

- **How many of all orders were required on the same day they were ordered?**
- **Which products were sold in a year that predates their model year?**
- **How many direct reports does each manager have? Just identify the manager by their staff ID.**

Typical join “shortcut”

- Often join tables on a common set of columns and then ask to only show that common column once.
 - ▶ Currently shown using the “from” clause, the “where” clause, and the choice of columns
- Modify the “from” clause to a “from ... join .. using (..)” clause:
 - ▶ “join” replacing the comma in the from clause
 - ▶ “using” lists the columns to have in common
 - ▶ The variant only keeps one copy of all the columns mentioned in the “using” part

Natural Join

- When your join is going to use all common columns, you can use the "natural join" operator
 - ▶ Equivalent to join..using (<all common column names>)

Example

- **Select course.course_id, name, person_id, fees
from course, registration
where course.course_id = registration.course_id;**

becomes

**Select *
from course join registration using (course_id);**

or

Select * from course natural join registration;

Different kinds of join operators

● Consider 4 join operators:

▶ Inner join

- Returns records that have matching values in both tables
- **Equi-join** – inner join where the join condition is based on equality between values in the common columns
- **Natural join** – inner join that restricts all same-named columns to match and produces one instance of the common columns

▶ Left join

- Return all records from the left table, and the matched records from the right table, adding NULL values when a match isn't present

▶ Right join

- Return all records from the right table, and the matched records from the left table, adding NULL values when a match isn't present

▶ Outer join

- Return all records when there is a match in either left or right table

Join examples

sample1

id	value	name
1	100	one
1	200	two
2	300	three

sample2

id	info	value
2	alpha	500
2	beta	300
4	gamma	400

```
[mysql> select * from sample1 join sample2 where sample1.id = sample2.id;
```

id	value	name	id	info	value
2	300	three	2	alpha	500
2	300	three	2	beta	300

Equi-join

```
[mysql> select * from sample1 natural join sample2;
```

id	value	name	info
2	300	three	beta

Natural join

Join examples

sample1

id	value	name
1	100	one
1	200	two
2	300	three

sample2

id	info	value
2	alpha	500
2	beta	300
4	gamma	400

```
mysql> select * from sample1 left join sample2 on sample1.id = sample2.id;
```

id	value	name	id	info	value
2	300	three	2	alpha	500
2	300	three	2	beta	300
1	100	one	NULL	NULL	NULL
1	200	two	NULL	NULL	NULL

Left join

```
mysql> select * from sample1 right join sample2 on sample1.id = sample2.id;
```

id	value	name	id	info	value
2	300	three	2	alpha	500
2	300	three	2	beta	300
NULL	NULL	NULL	4	gamma	400

Right join

Join conditions

- Outer joins often invoke “is null” or “is not null” in the where condition to filter the results

sample1			sample2		
id	value	name	id	info	value
1	100	one	2	alpha	500
1	200	two	2	beta	300
2	300	three	4	gamma	400

- What does the following produce?
Select sample2.id from sample1 right join sample2 on sample1.id = sample2.id where sample1.id is null;

- ▶ All ids in sample2 not in sample1

```
[mysql> select sample2.id from sample1 right join sample2  
[    -> on sample1.id = sample2.id where sample1.id is null;
```

id
4

Query Execution

- **The DBMS creates an execution plan from your SQL select statement**
 - ▶ **Identifies the order in which to do evaluations**
 - From: combine small tables first
 - Where: apply the most restrictive conditions first
 - ▶ **Query optimizer can re-order elements of the query to increase its performance**
 - Estimate the size of the tables to be combined to help manage the total work
- **How you specify your query can influence performance.**

Costs for select statements

- **Column selection**
 - ▶ Identifies what to keep as you process. Low cost, unless you're using transformations
- **From clause**
 - ▶ Generates combinations of records. High cost if you're generating many records that you will just throw away
- **Where clause**
 - ▶ Does winnowing as you process. Can have complex sets of conditions to evaluate. Medium cost
- **Group by clause**
 - ▶ Need the final data to make this work. Throwing away most of the generated records to create summaries.
- **Order by clause**
 - ▶ Need the final data to execute. Cost relative to output size, not generated records size.

Helping performance

- 1. Minimize the size and number of table combinations in the "from" clause
 - ▶ How??? We need a new tool -> subqueries
- 2. include restrictive "where" elements if you can

Subqueries

- The output of an SQL statement is a table.
- Use that output table in the place of any other table in a query.
 - ▶ Enclose the subquery in parentheses
 - ▶ Need to use the “as” keyword to give the output of the subquery a name
- Use a subquery to reduce the size or number of tables to combine in the “from” clause
- A subquery can have its own subquery

Subquery example

- **Select ***
from person join registration using (person_id)
where person_id = 3;

versus

select *
from
(select * from person where person_id = 3) as interested
join registration using (person_id);

Is there a difference?

Subquery example

● **select ***
from person as p, course as c, registration as r
where p.person_id = r.person_id
and c.course_id = r.course_id
and p.name like "A%";

*How many rows
does each query
create?*

versus

select * from
(
(select * from person where name like "A%") as p
join registration using (person_id)
) join course as c using (course_id);

Subqueries

- Can also appear in the “where” clause

- ▶ Extract one number for a comparison

```
select name, salary from person
  where salary >= (select avg(salary) from person);
```


- ▶ Extract a set of values for an “in” statement (waiting for mysql to catch up to this functionality)

```
select name, salary from person
  where salary in
    (select distinct salary from person order by salary desc limit
2);
```


Find everyone within 1 standard deviation of the average age – shorter version from class

- **Select name, age from person where age between
(select avg(age) – std(age) from person)
and
(select avg(age) + std(age) from person);**

Bicycle store database



Database from
<http://www.sqlservertutorial.net/load-sample-database/>
 ER diagram reverse engineered by MySQL Workbench

Question to try...

- Which products are out of stock in one or more stores, but company as a whole has some somewhere? Which store is out of stock?